

# Simplifying Image Processing Using Dimensionality Reduction

---



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Dictionary learning for dimensionality reduction of image data**

**Feature detection using convolutional layers**

**Autoencoders for dimensionality reduction**

# Dictionary Learning and Feature Extraction

---

# Dictionary Learning

Representation learning method to find a sparse representation of input data, often used in denoising of images.

# Dictionary Learning for Image Denoising

## Corpus of clean images

Should be free of noise

Feed into Dictionary Learner

## Choose transformation algorithm

Implement numeric procedures

Several choices - Orthogonal Matching Pursuit (OMP), LARS

## Perform Dictionary Learning

Express clean images in terms of atoms (sparse codings)

Each image is expressed as linear combination of atoms

## Use of Image Denoising

Apply to new, noisy images

Reconstruct and express using atoms

# Dictionary Learning

## Several choices of solver

- Thresholding: Fast but inaccurate
- Orthogonal Matching Pursuit (OMP):  
Most accurate, unbiased
- Least-angle Regression
- Lasso Regression

Demo

**Dictionary learning to learn sparse  
codings**

# Convolution Kernels for Feature Detection

---



“Sometimes the mind can see what is invisible to the eye”

# Viewing an Image



**All neurons in the eye don't see the entire image**

# Viewing an Image



**Each neuron has its own local receptive field**

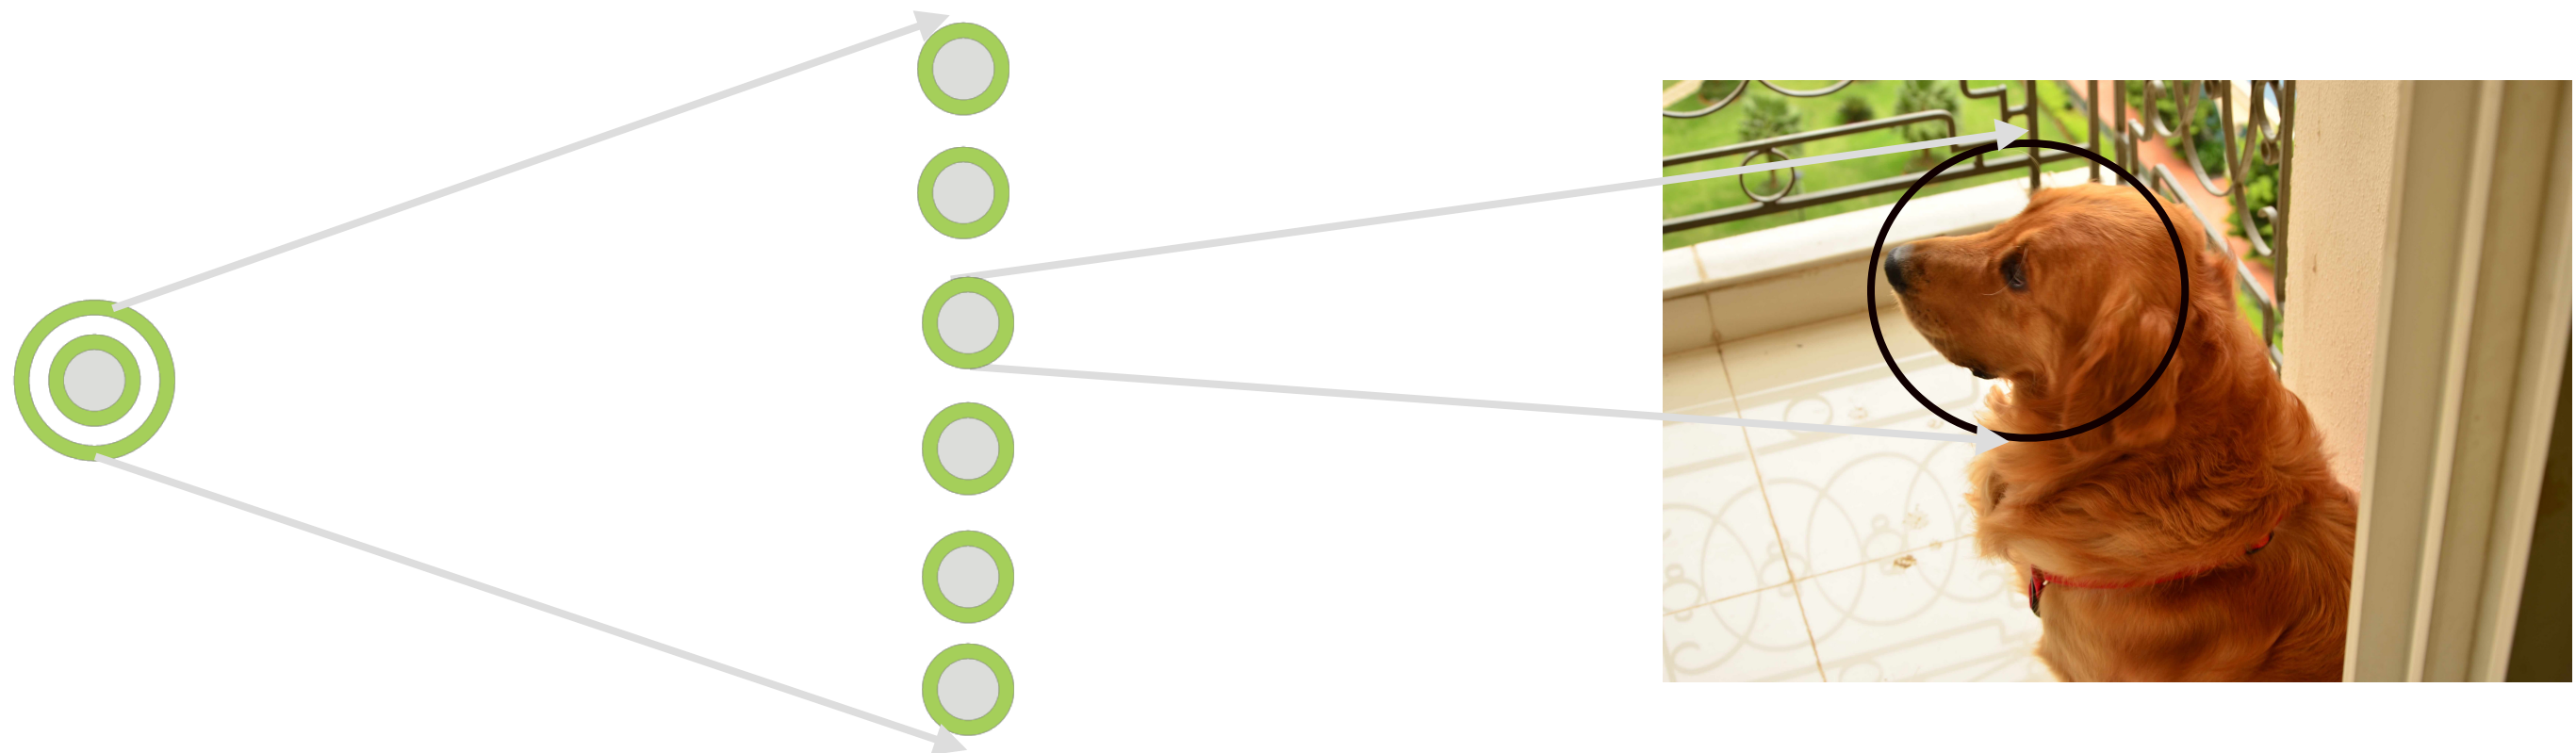
# Viewing an Image



**It reacts only to visual stimuli  
located in its receptive field**



# Viewing an Image



Some neurons react to more complex patterns  
that are **combinations** of lower level patterns

# Convolution

In this context, a sliding window function applied to a matrix

# Convolution

In this context, a sliding window function applied to  
a matrix



e.g. a matrix of pixels  
representing an image

# Convolution

In this context, a sliding window **function** applied to a matrix

Often called a kernel or filter





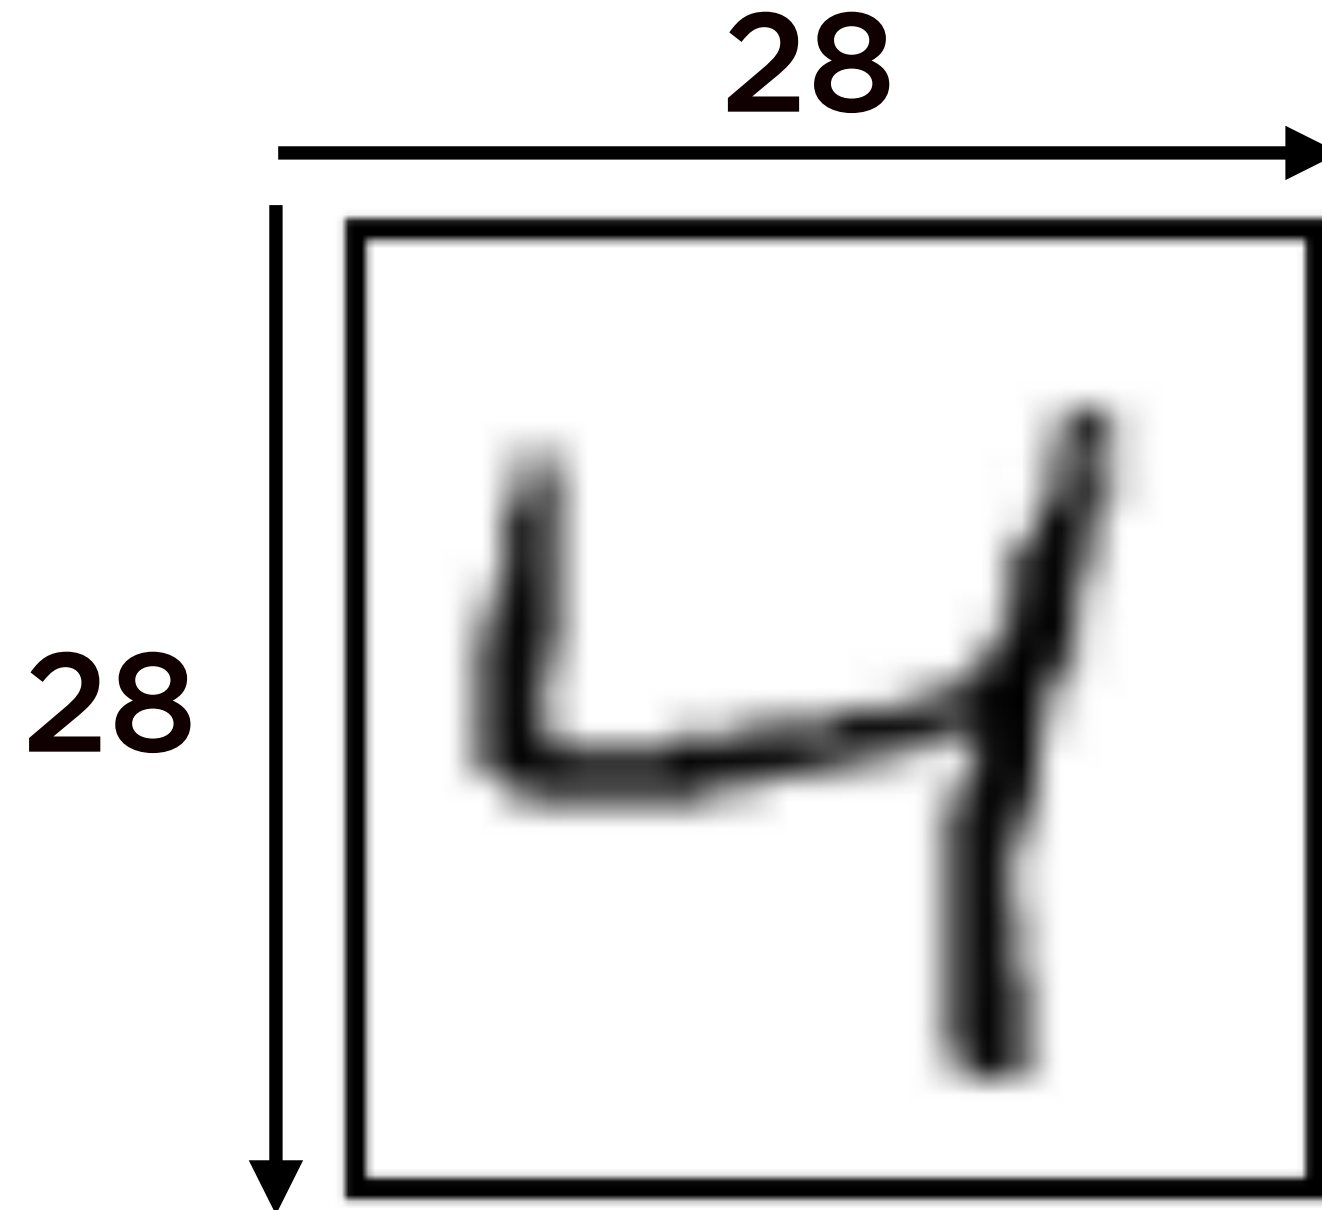
# Convolution

In this context, a sliding window function applied to a matrix



Kernel is applied element-wise  
in sliding-window fashion

# Representing Images as Matrices



**= 784 pixels**

# Representing Images as Matrices

6

6

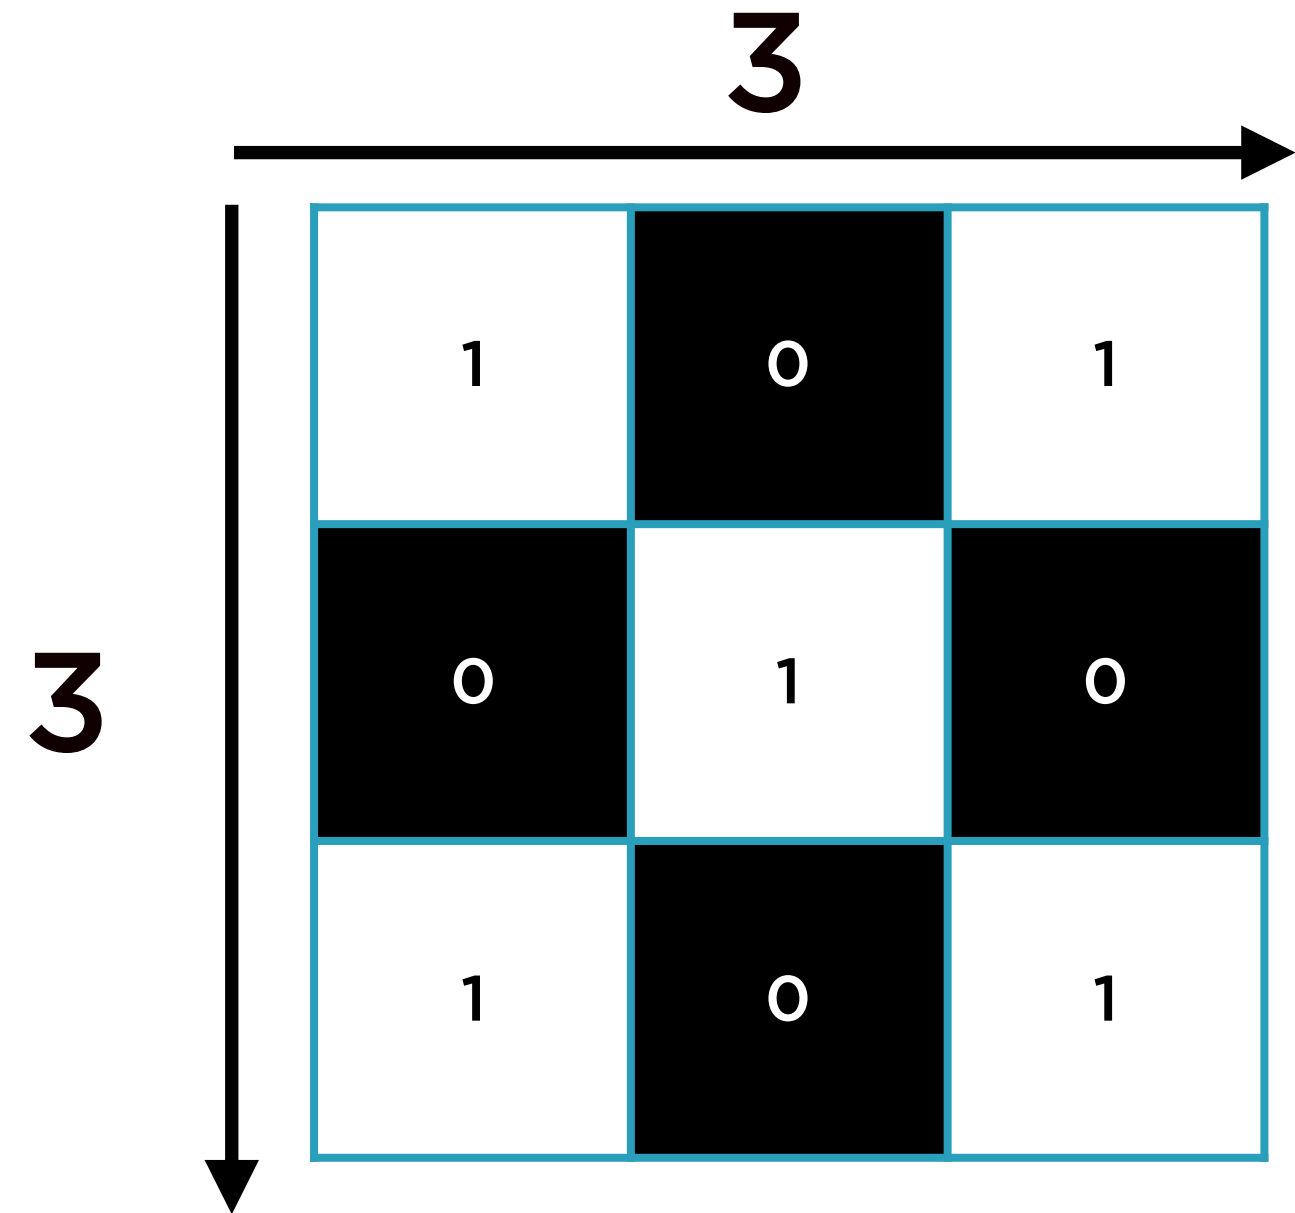
0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

= 36 pixels

# Representing Images

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

**Matrix**

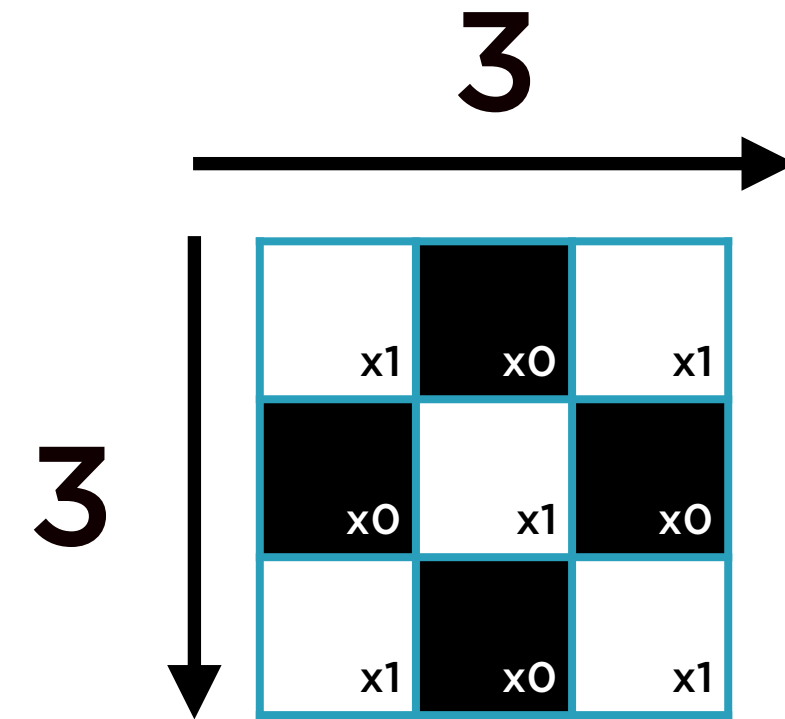


**Kernel**

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

**Matrix**



**Kernel**

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



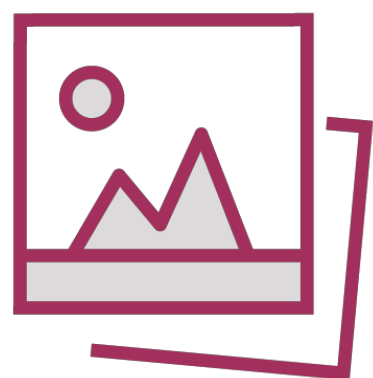
x1	x0	x1
x0	x1	x0
x1	x0	x1

4

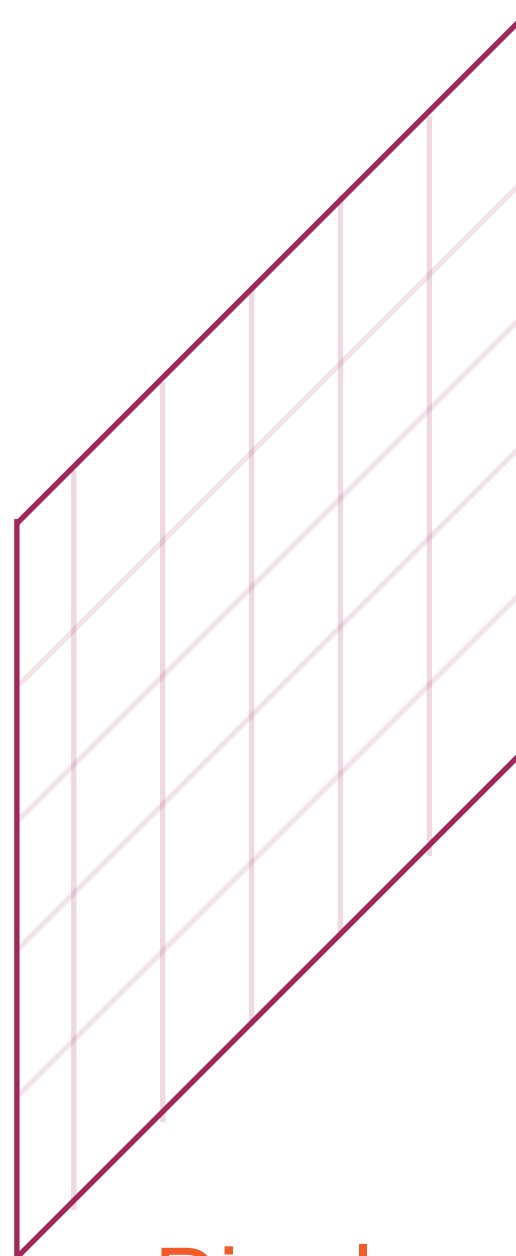
4			
→			
1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution  
Result

# Feature Maps

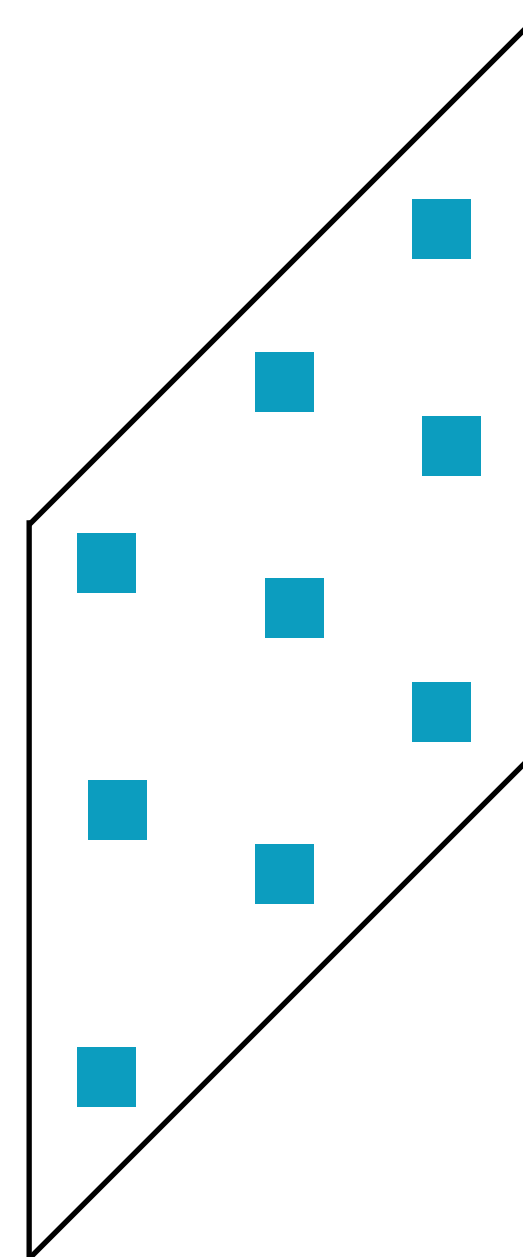


Image



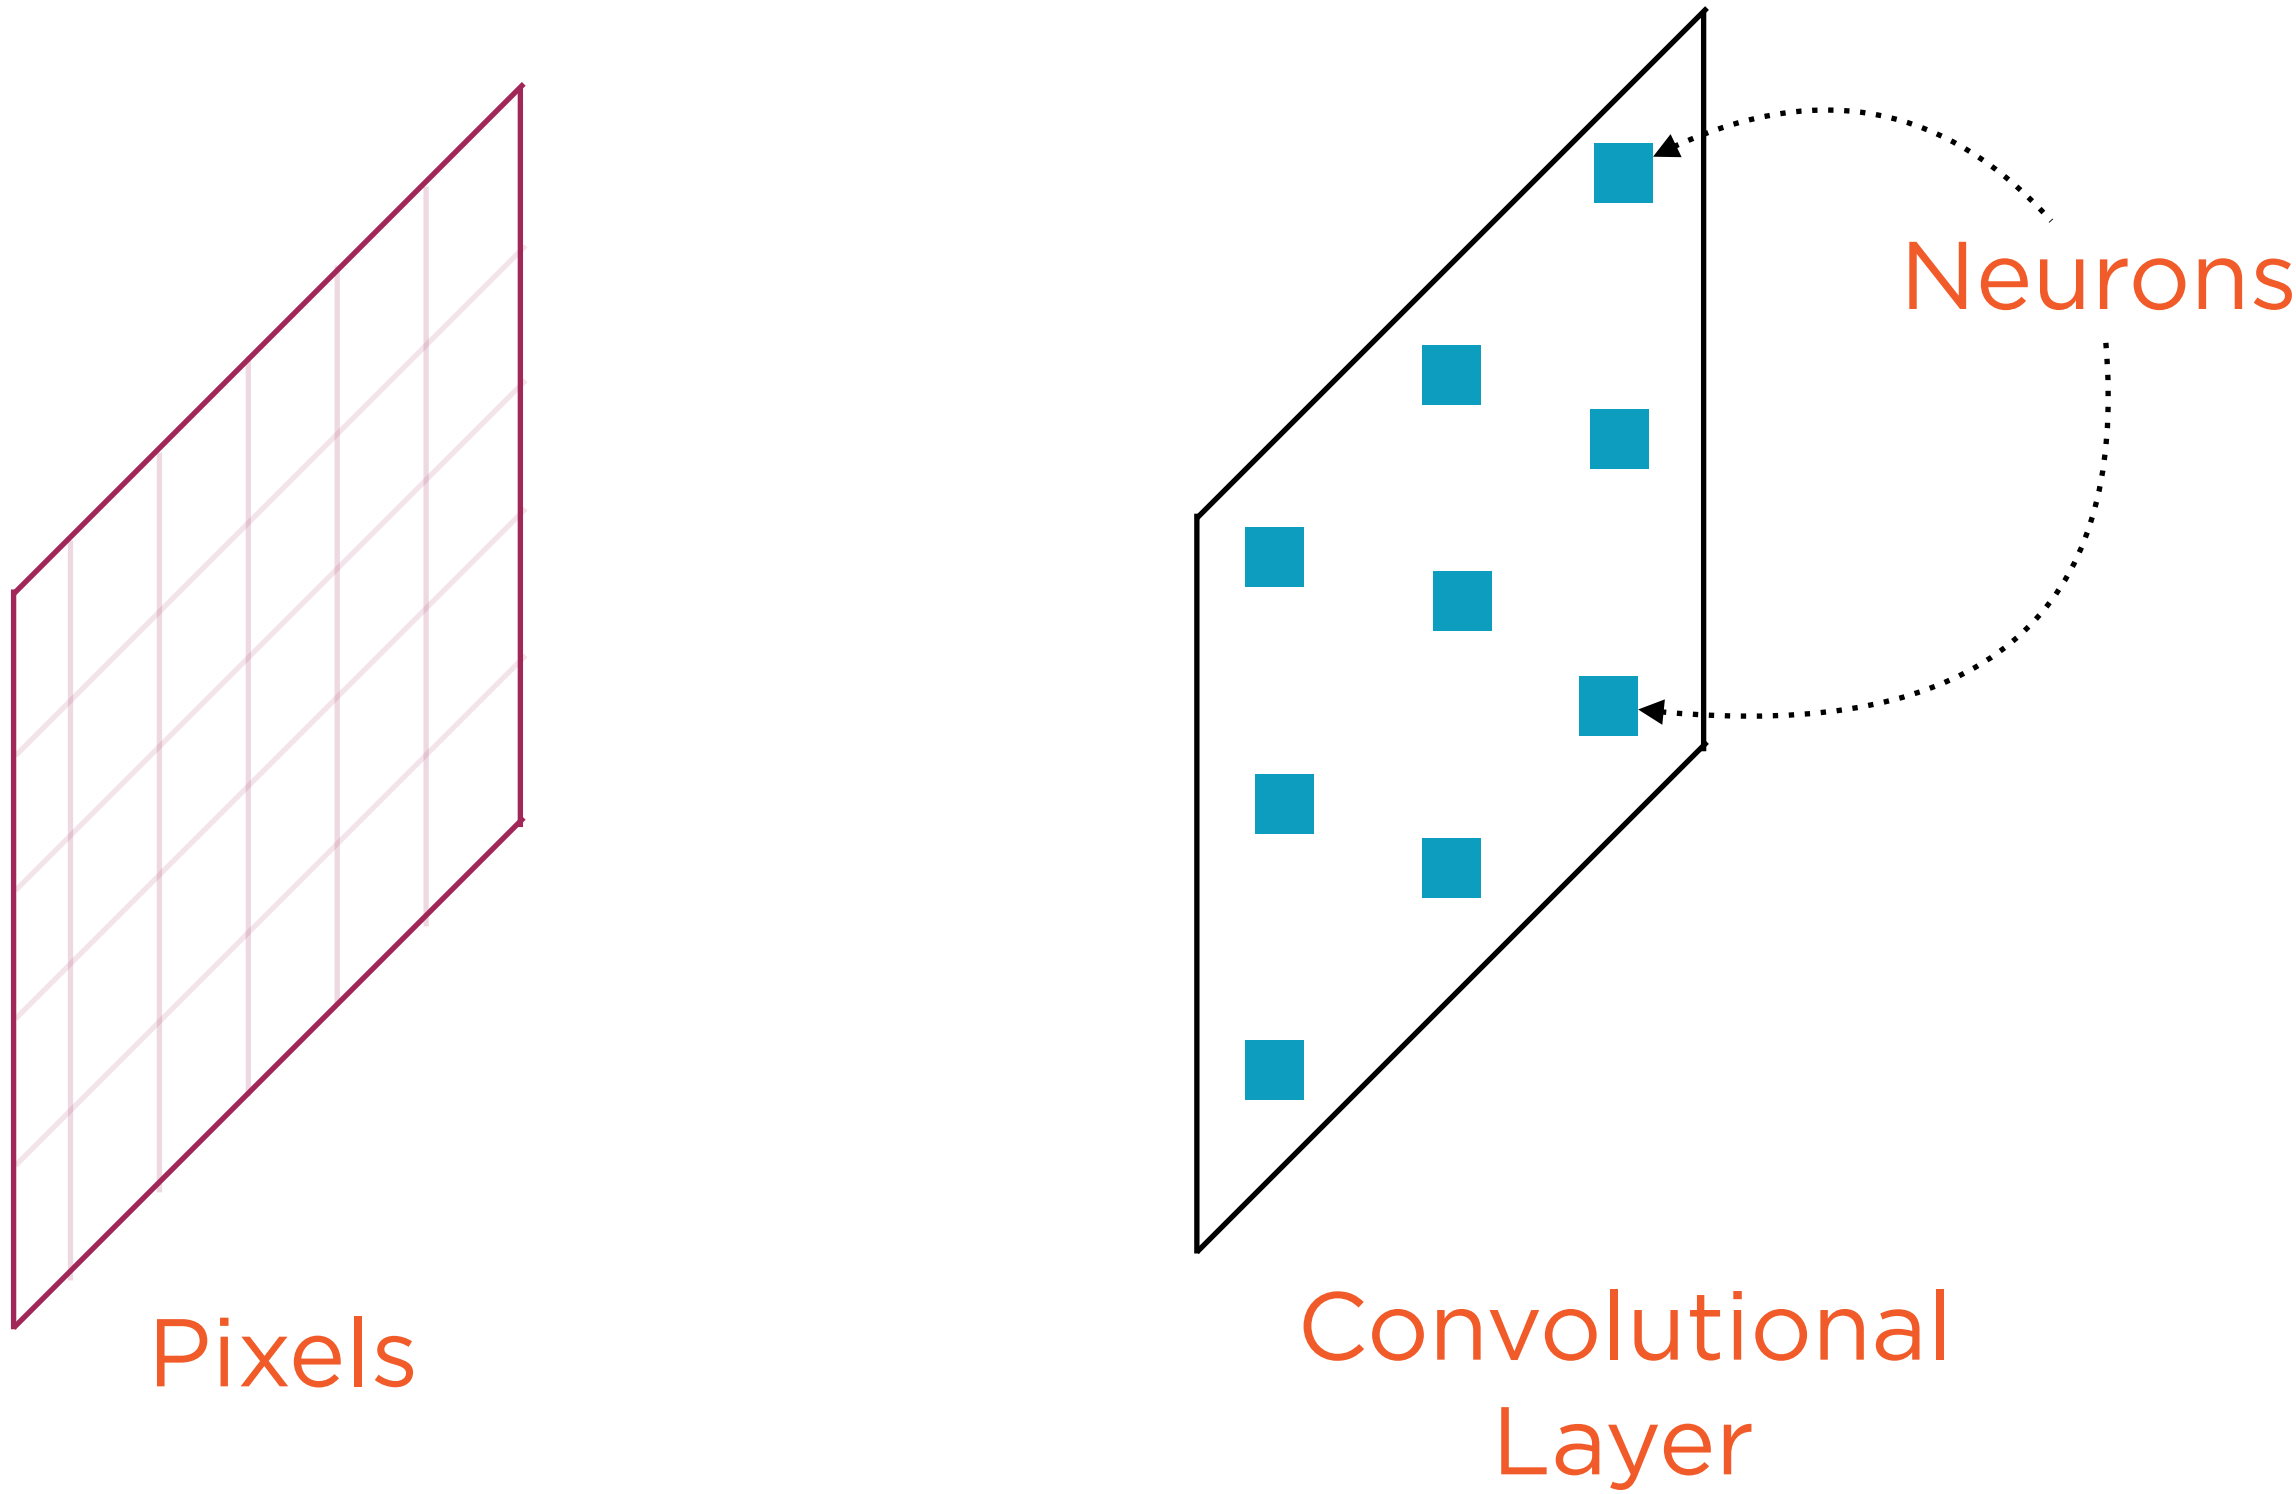
Pixels

x1	x0	x1
x0	x1	x0
x1	x0	x1



Feature  
Map

# Feature Maps

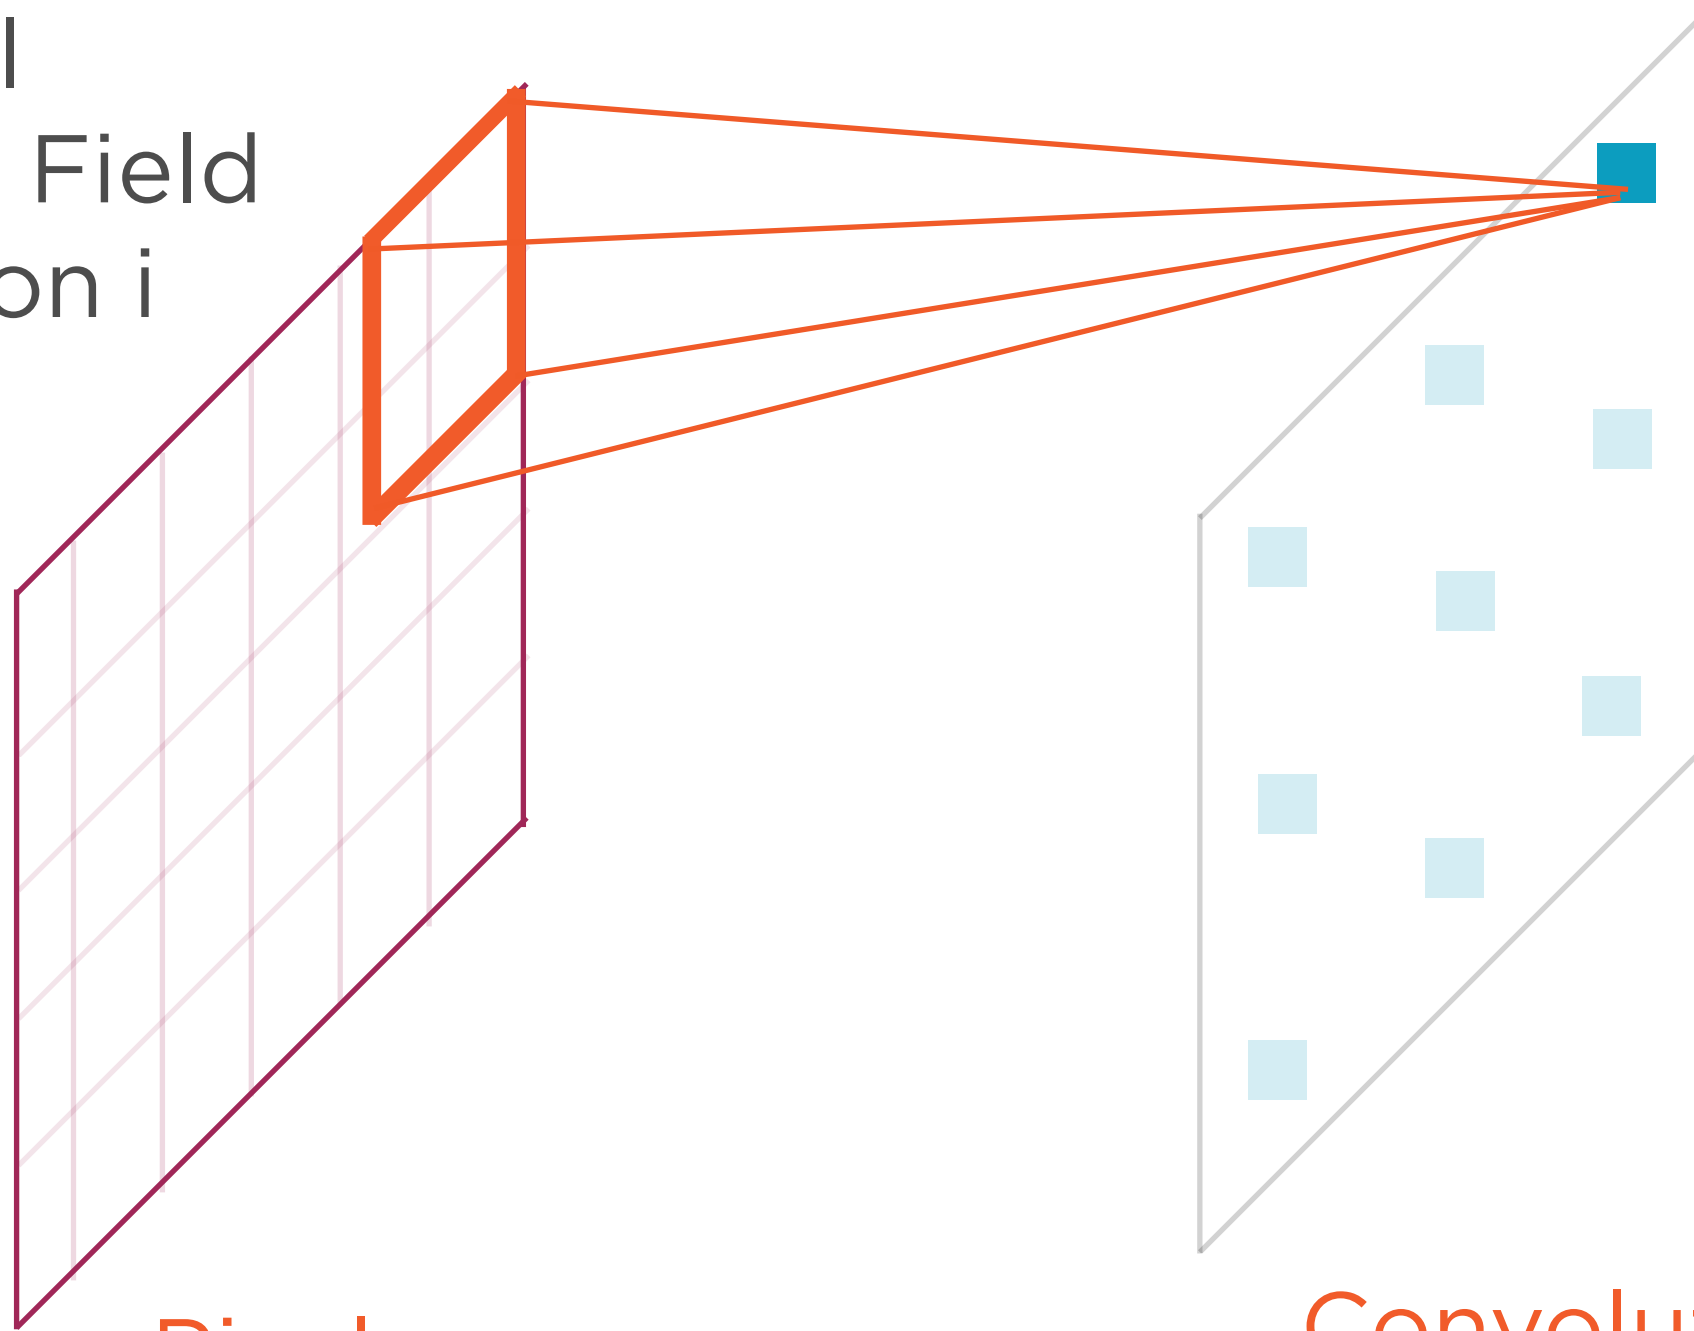




# Feature Maps

Local  
Receptive Field  
of Neuron  $i$

Neuron  $i$



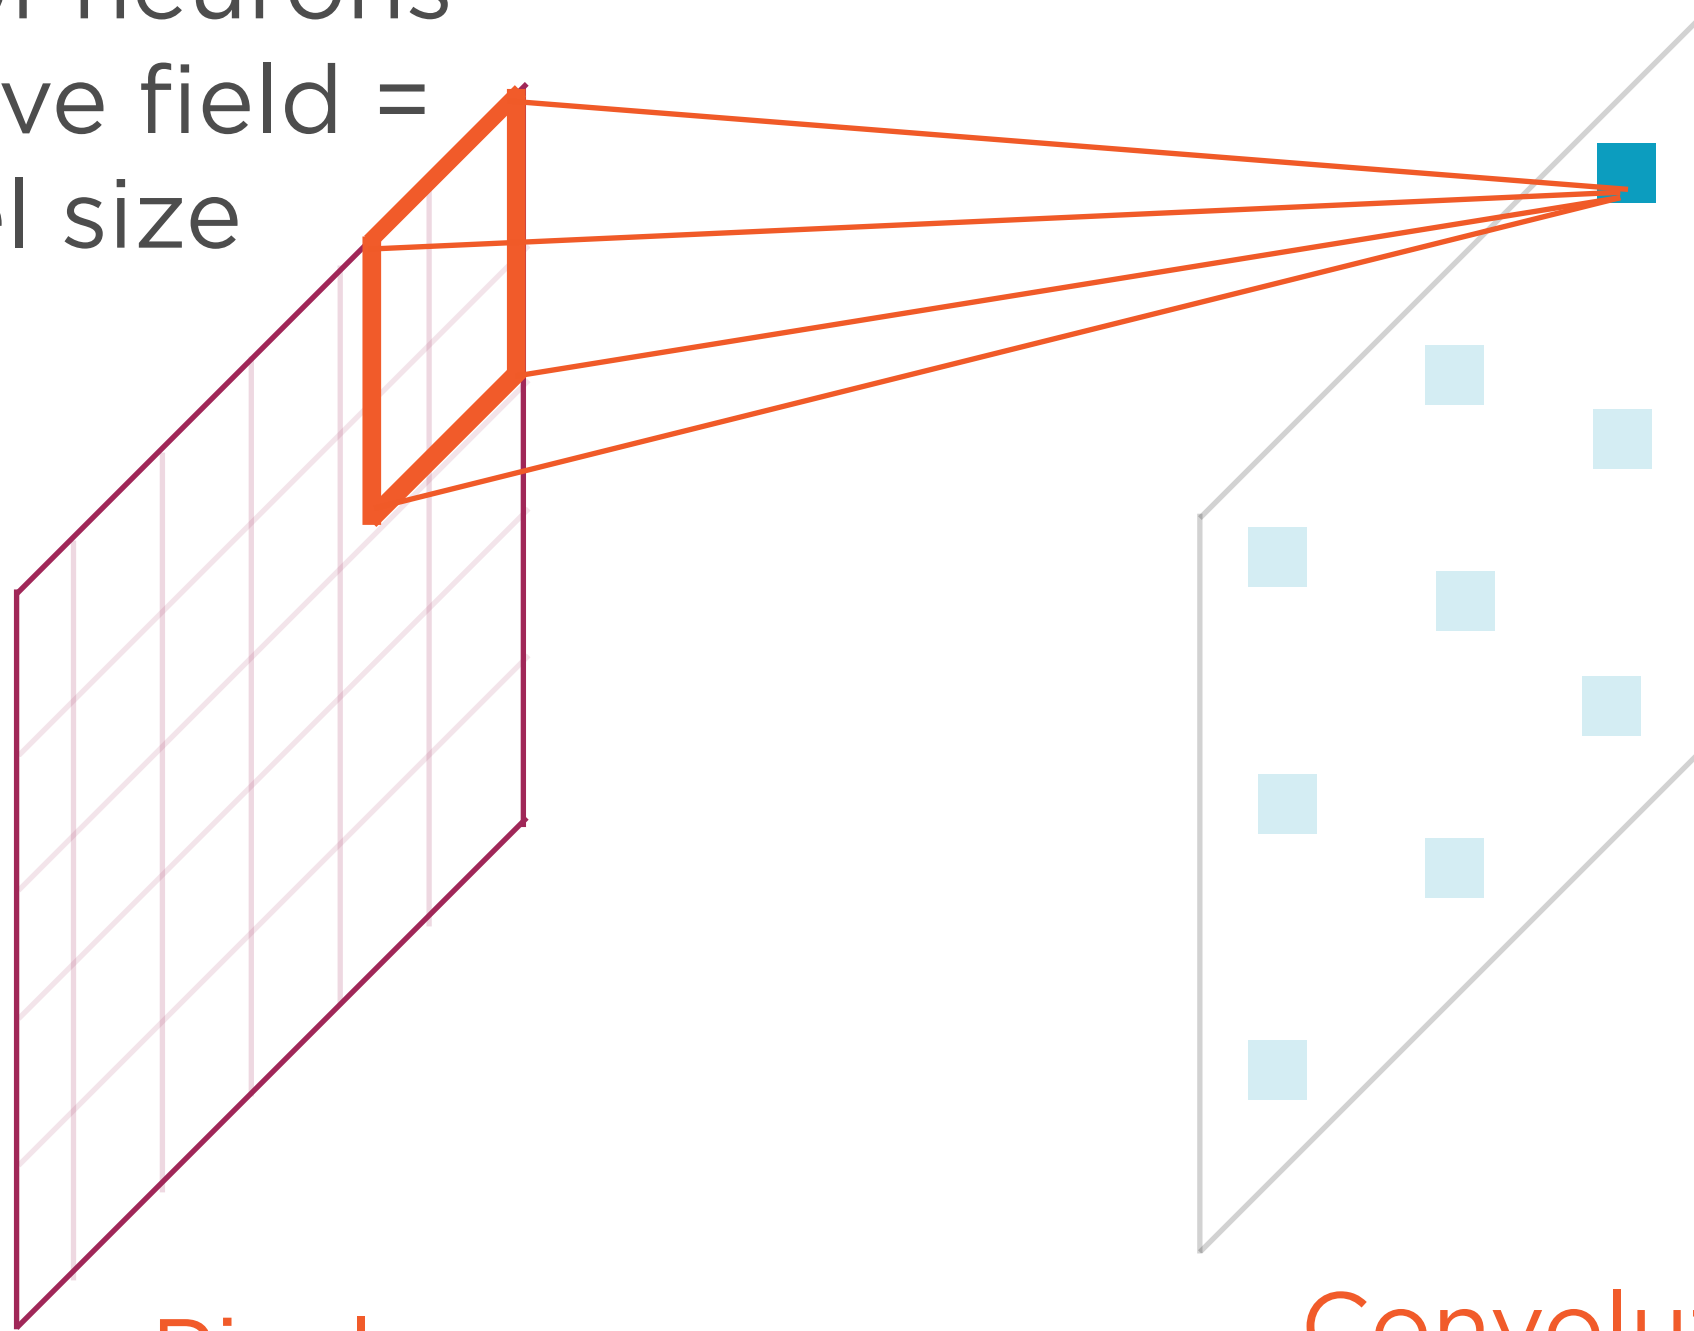
Pixels

Convolutional  
Layer

Number of neurons  
in receptive field =  
kernel size

Feature Maps

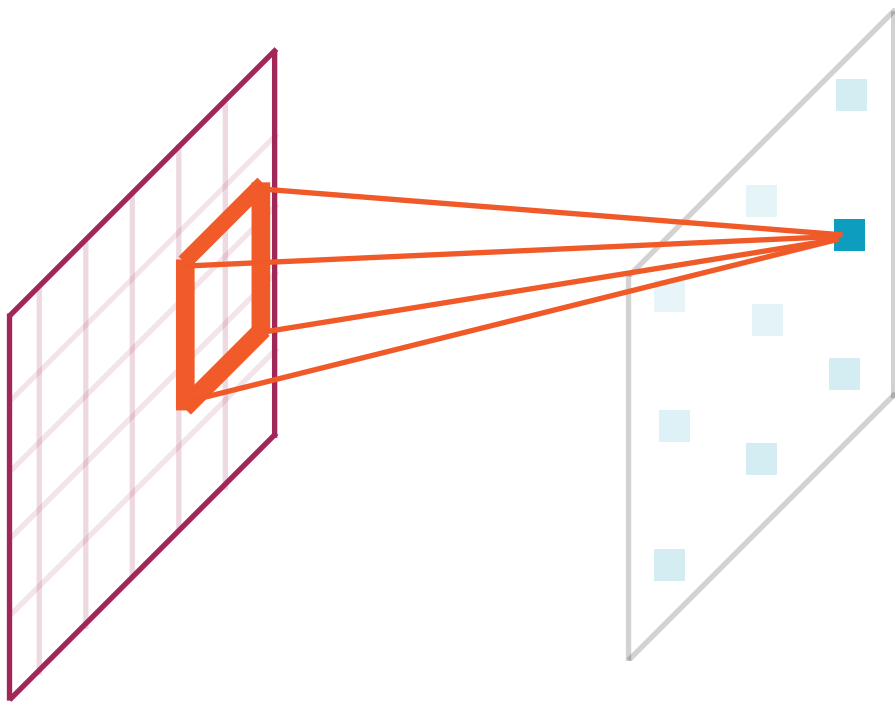
Neuron i



Pixels

Convolutional  
Layer

# Feature Maps

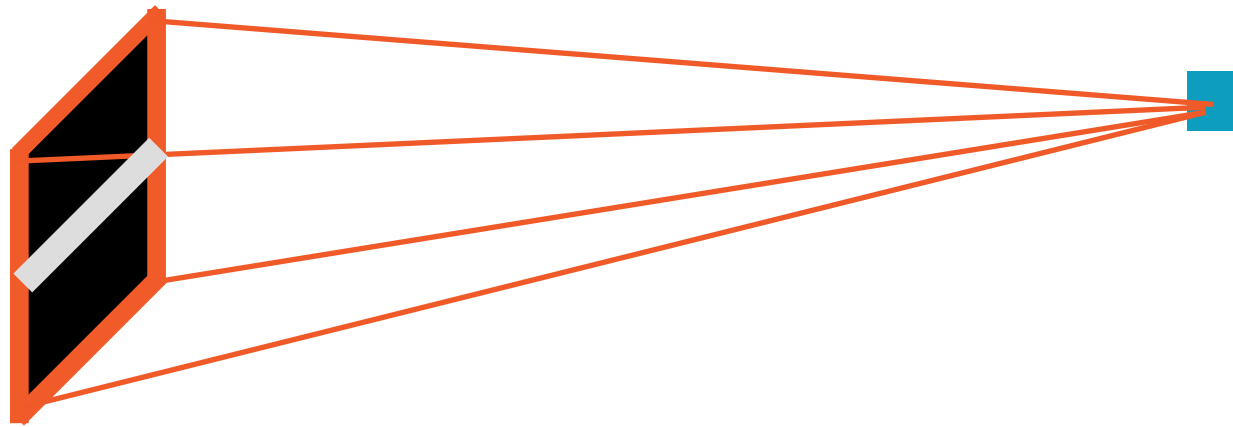


The parameters of all neurons in a feature map are collectively called the filter

Why filter?

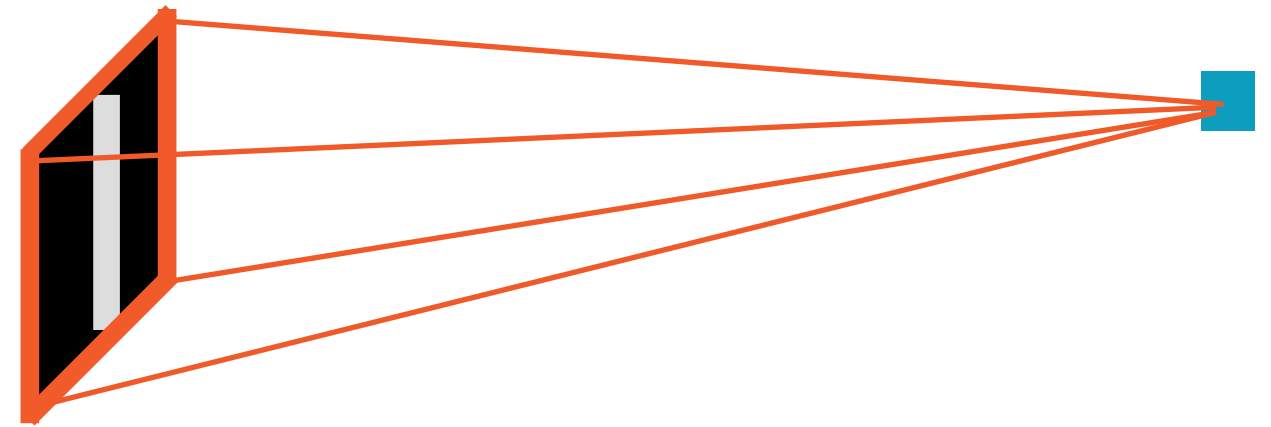
Because weights **highlight (filter)** specific patterns from the input pixels

# Filters



## Horizontal Filter

Neuron will detect  
horizontal lines in input



## Vertical Filter

Neuron will detect  
vertical lines in input

Demo

**Convolutional kernels for feature  
detection**

# Autoencoding

---

# Choosing Autoencoders

## Use Case

**Extremely complex feature  
vectors**

**Images, video, documents**

**Pre-processing before using in  
neural networks**

## Possible Solution

**Autoencoders**

$$y = f(x)$$

---

# Supervised Machine Learning

**Most machine learning algorithms seek to “learn” the function  $f$  that links the features and the labels**



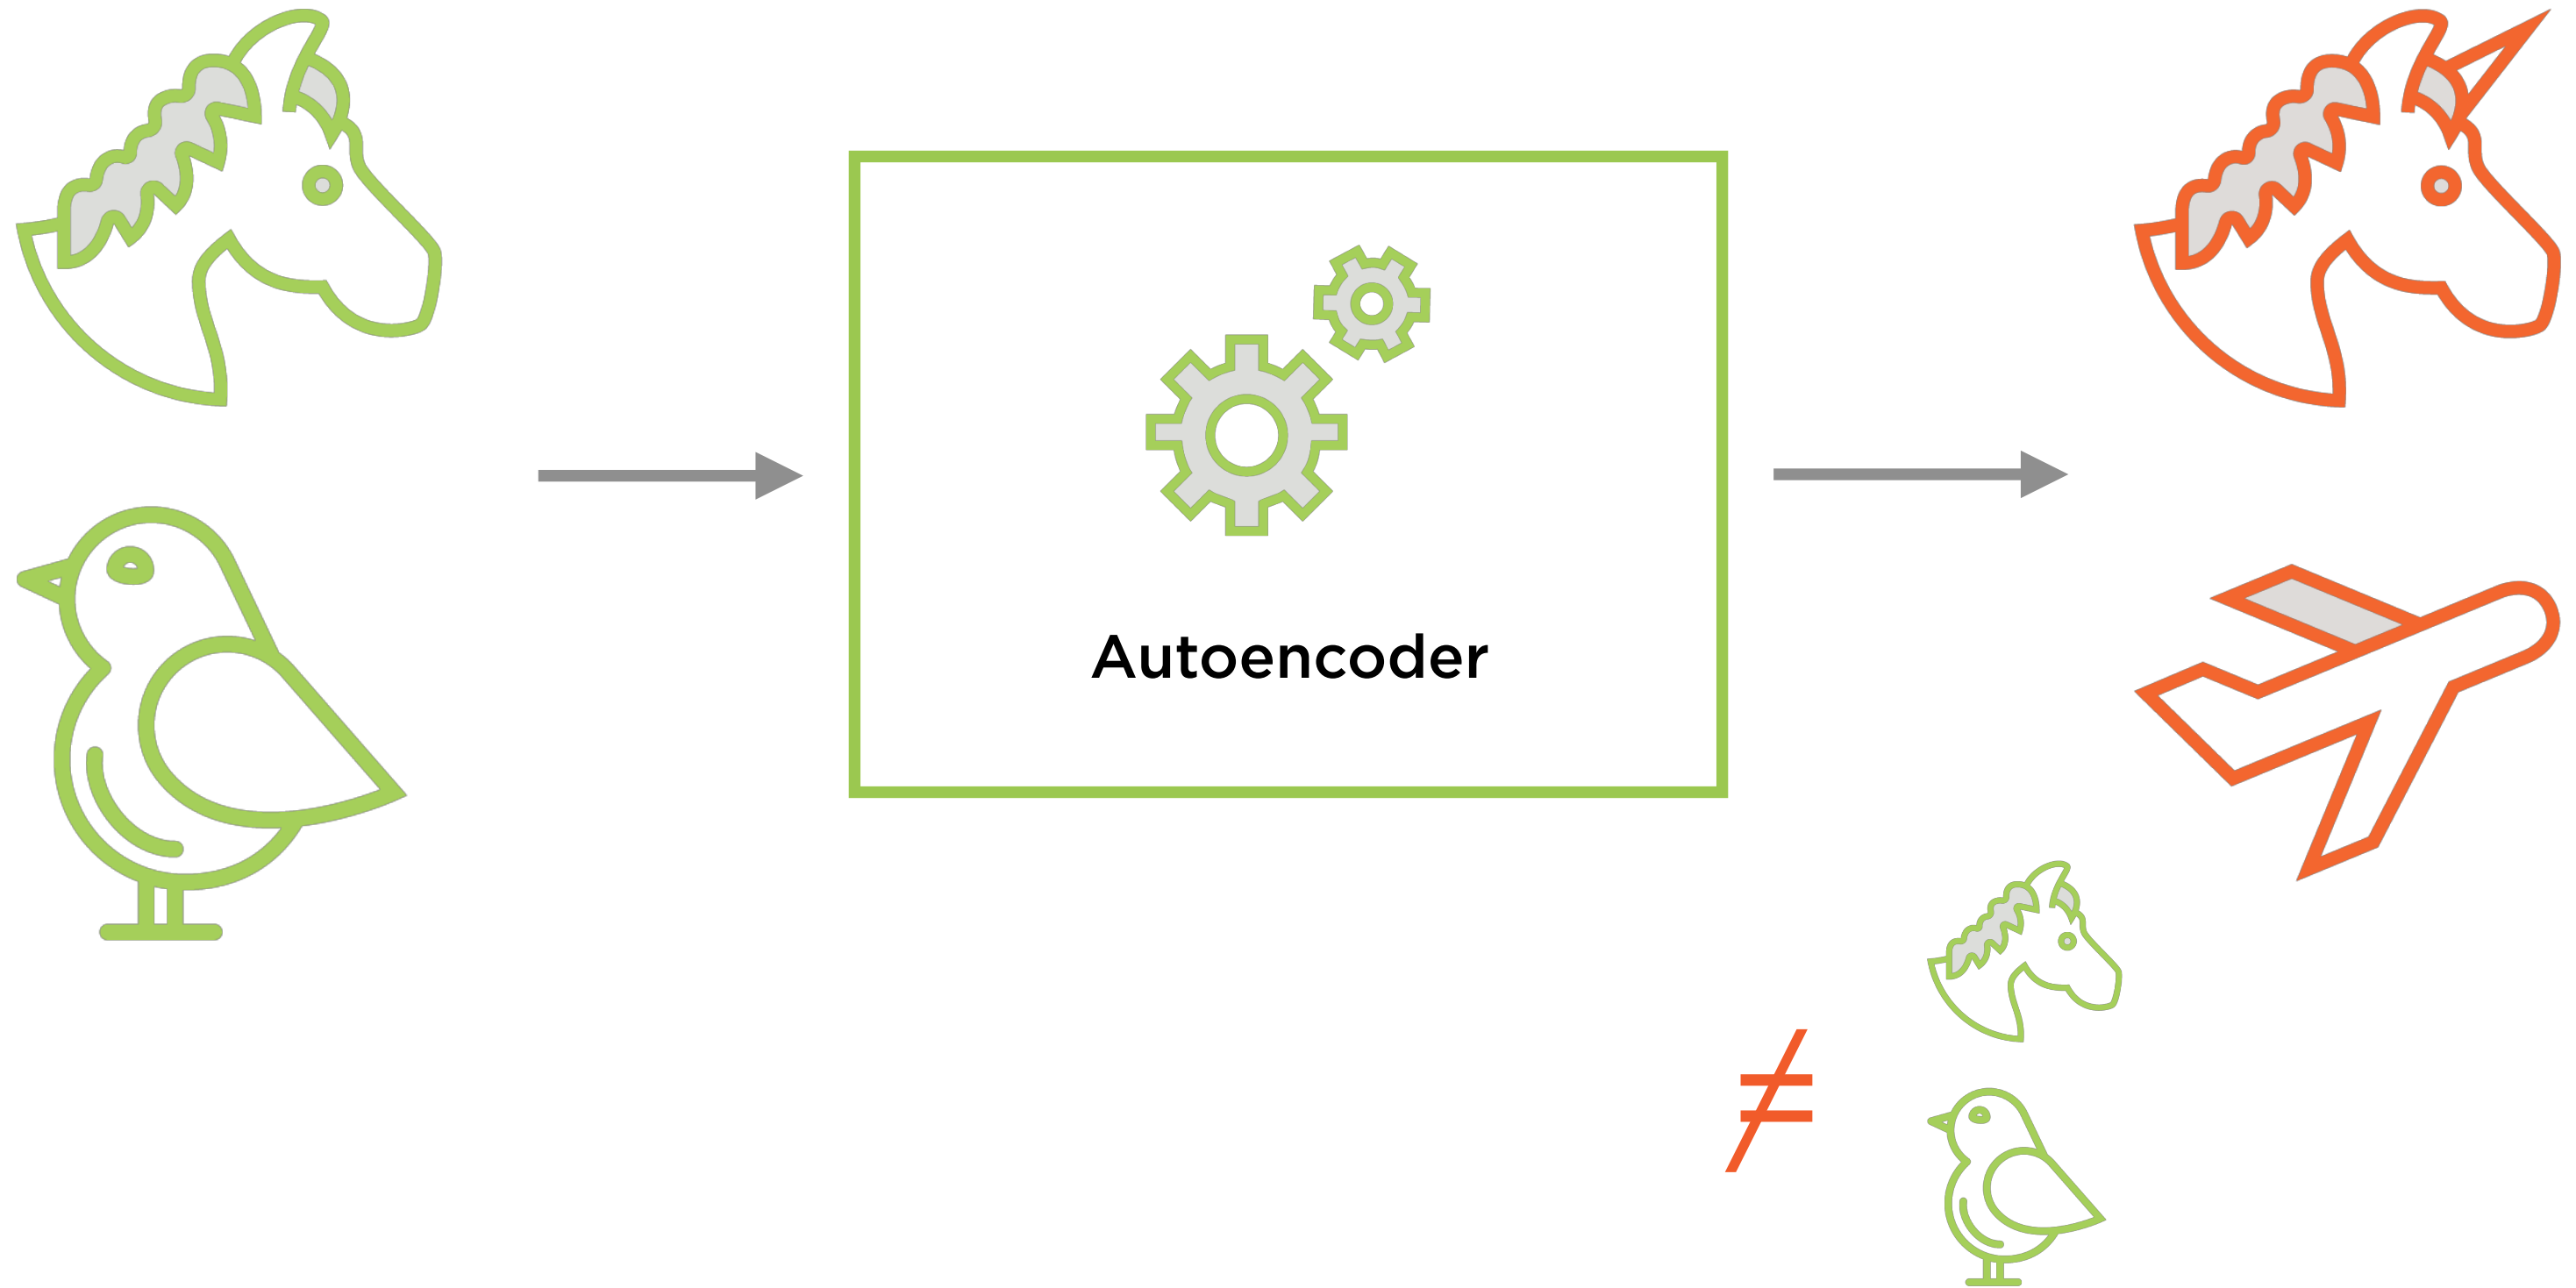
$$\mathbf{x} = \mathbf{f}(\mathbf{x})$$

---

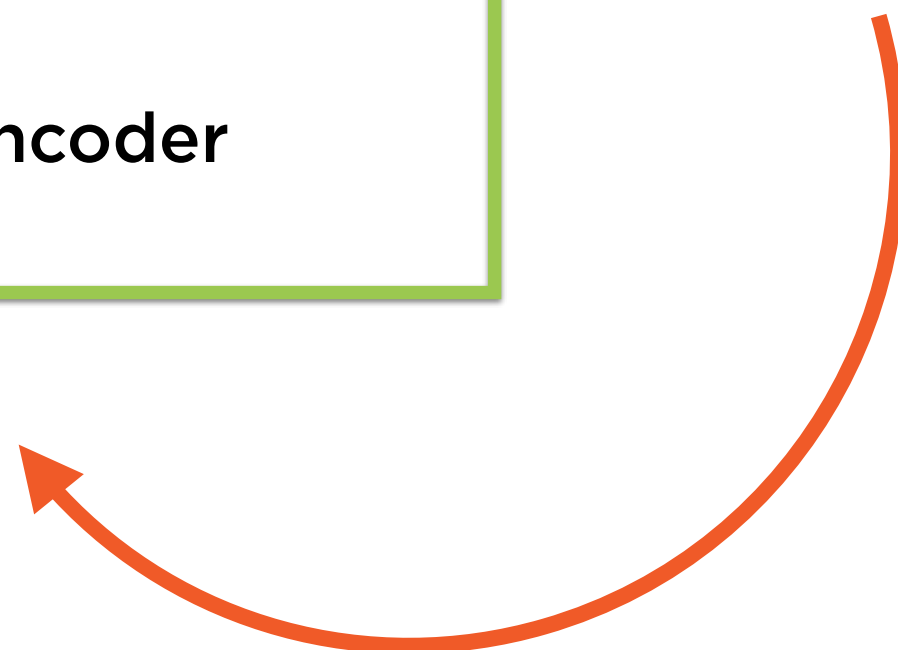
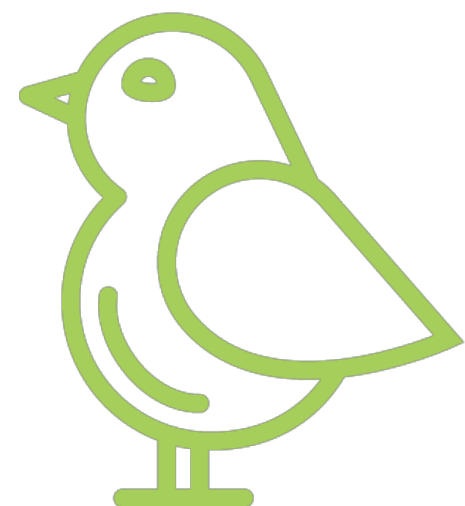
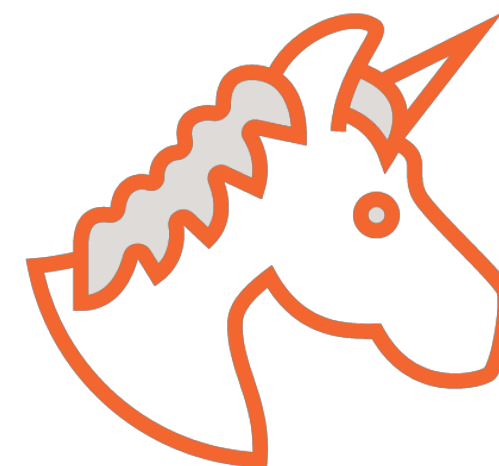
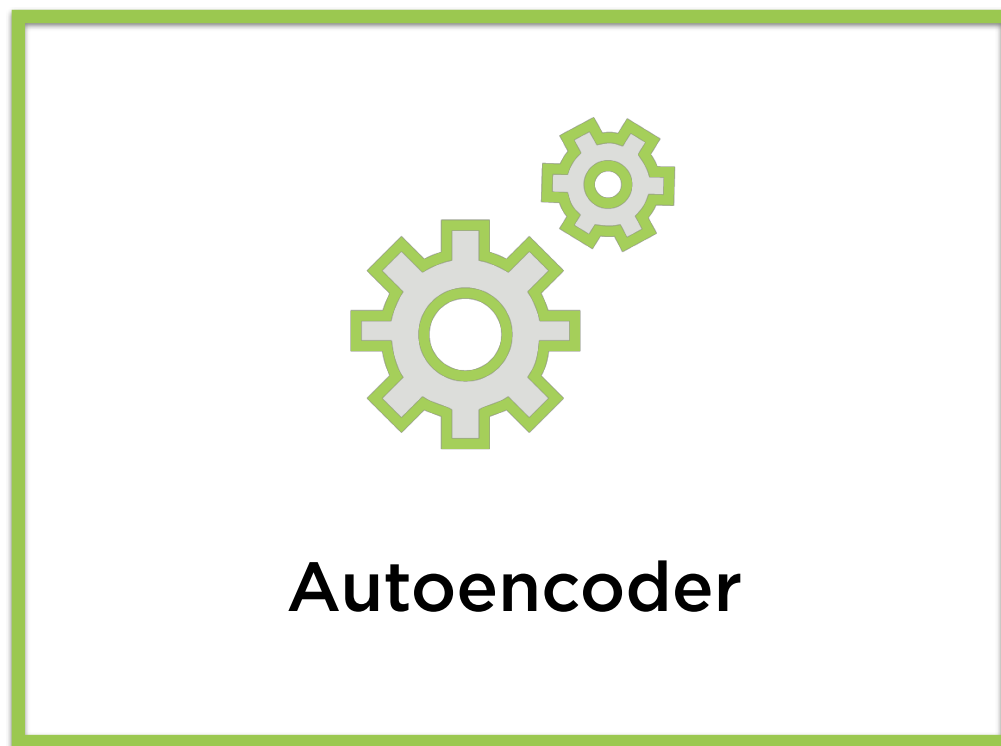
# Autoencoders Learn the Input!

**The process is inherently unsupervised, but cleverly uses the input itself to train an algorithm**

# Autoencoder

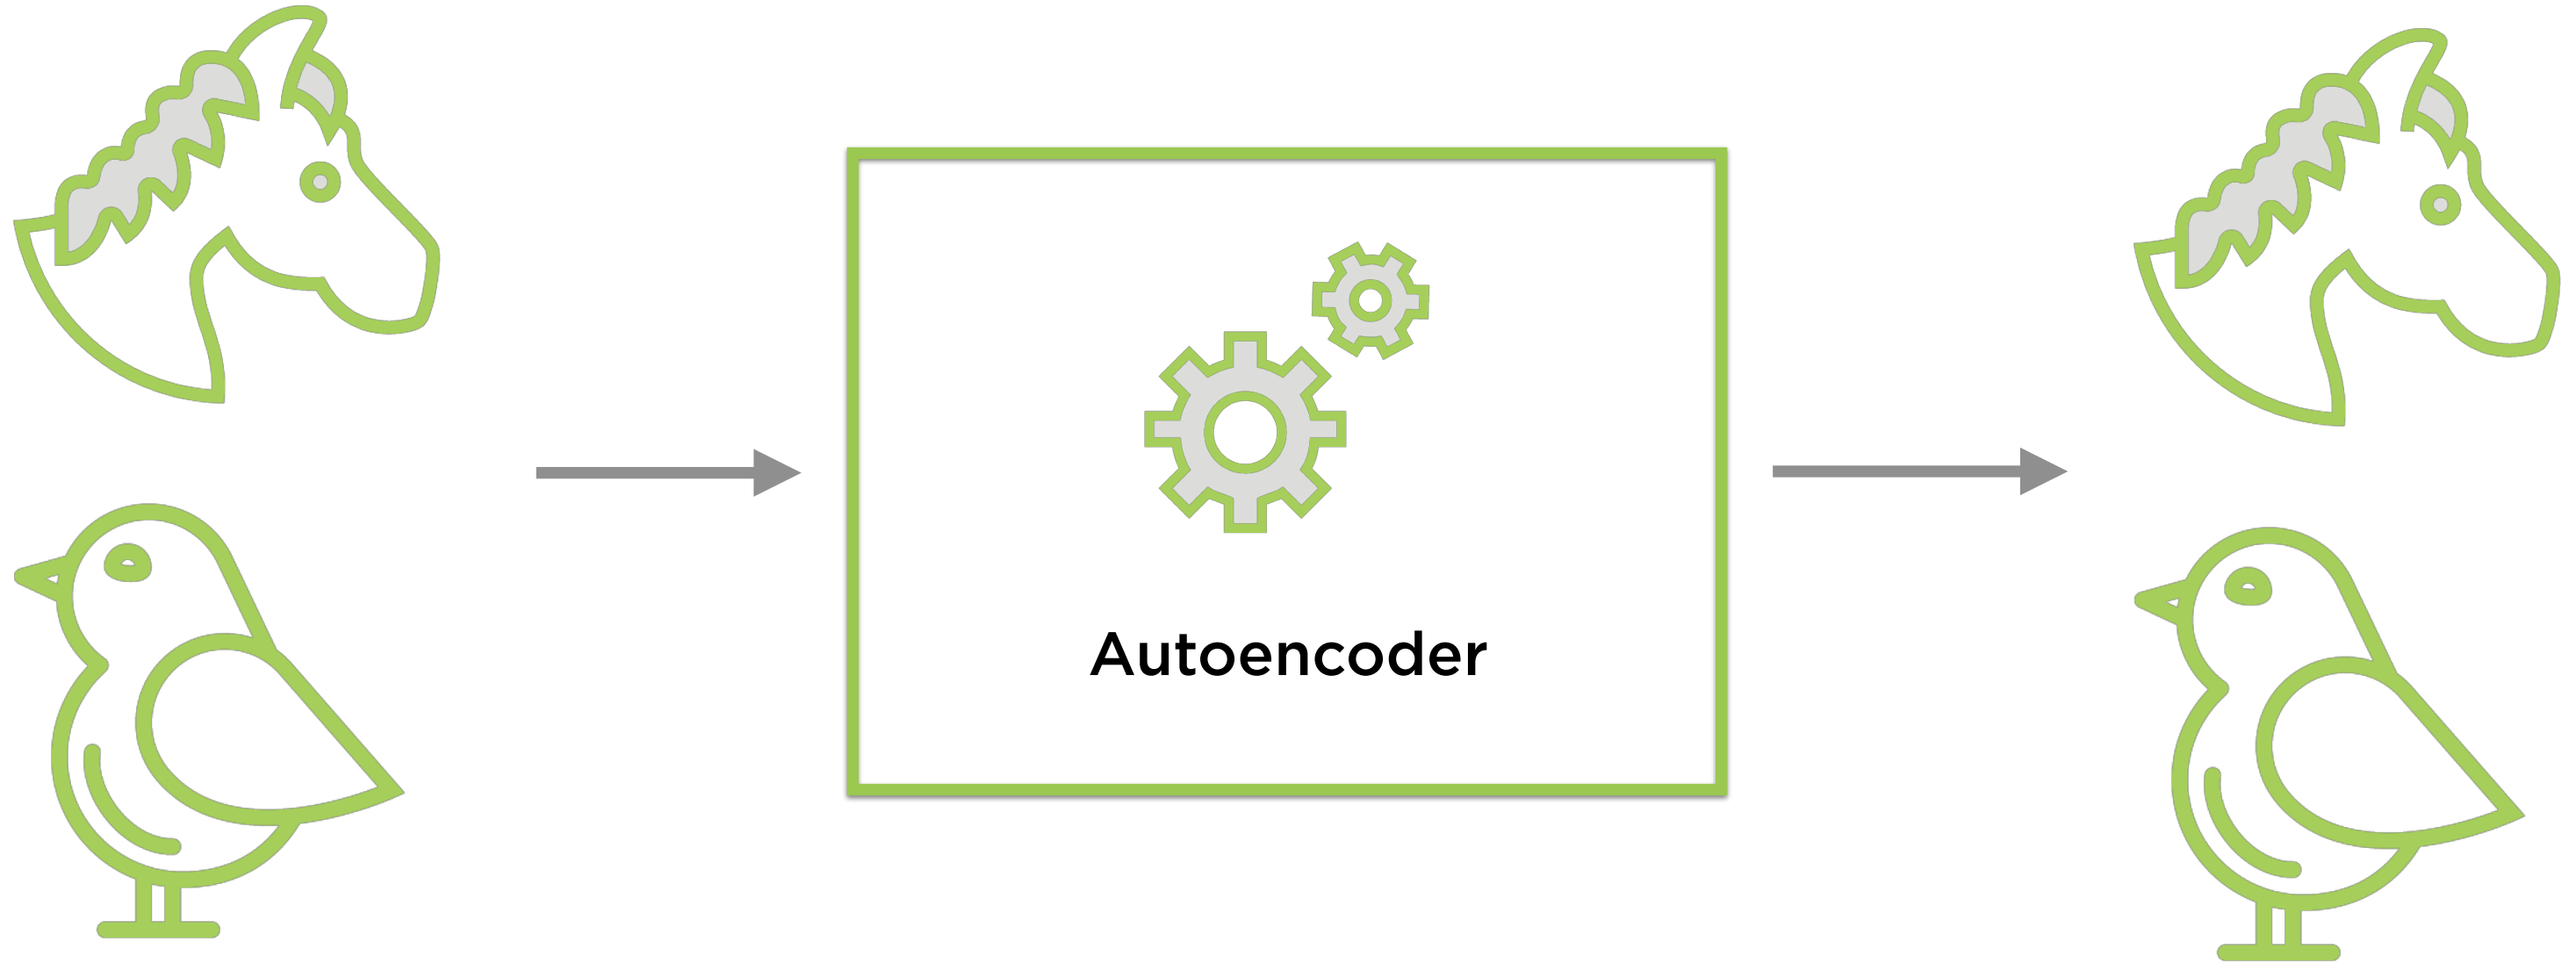


# Autoencoder



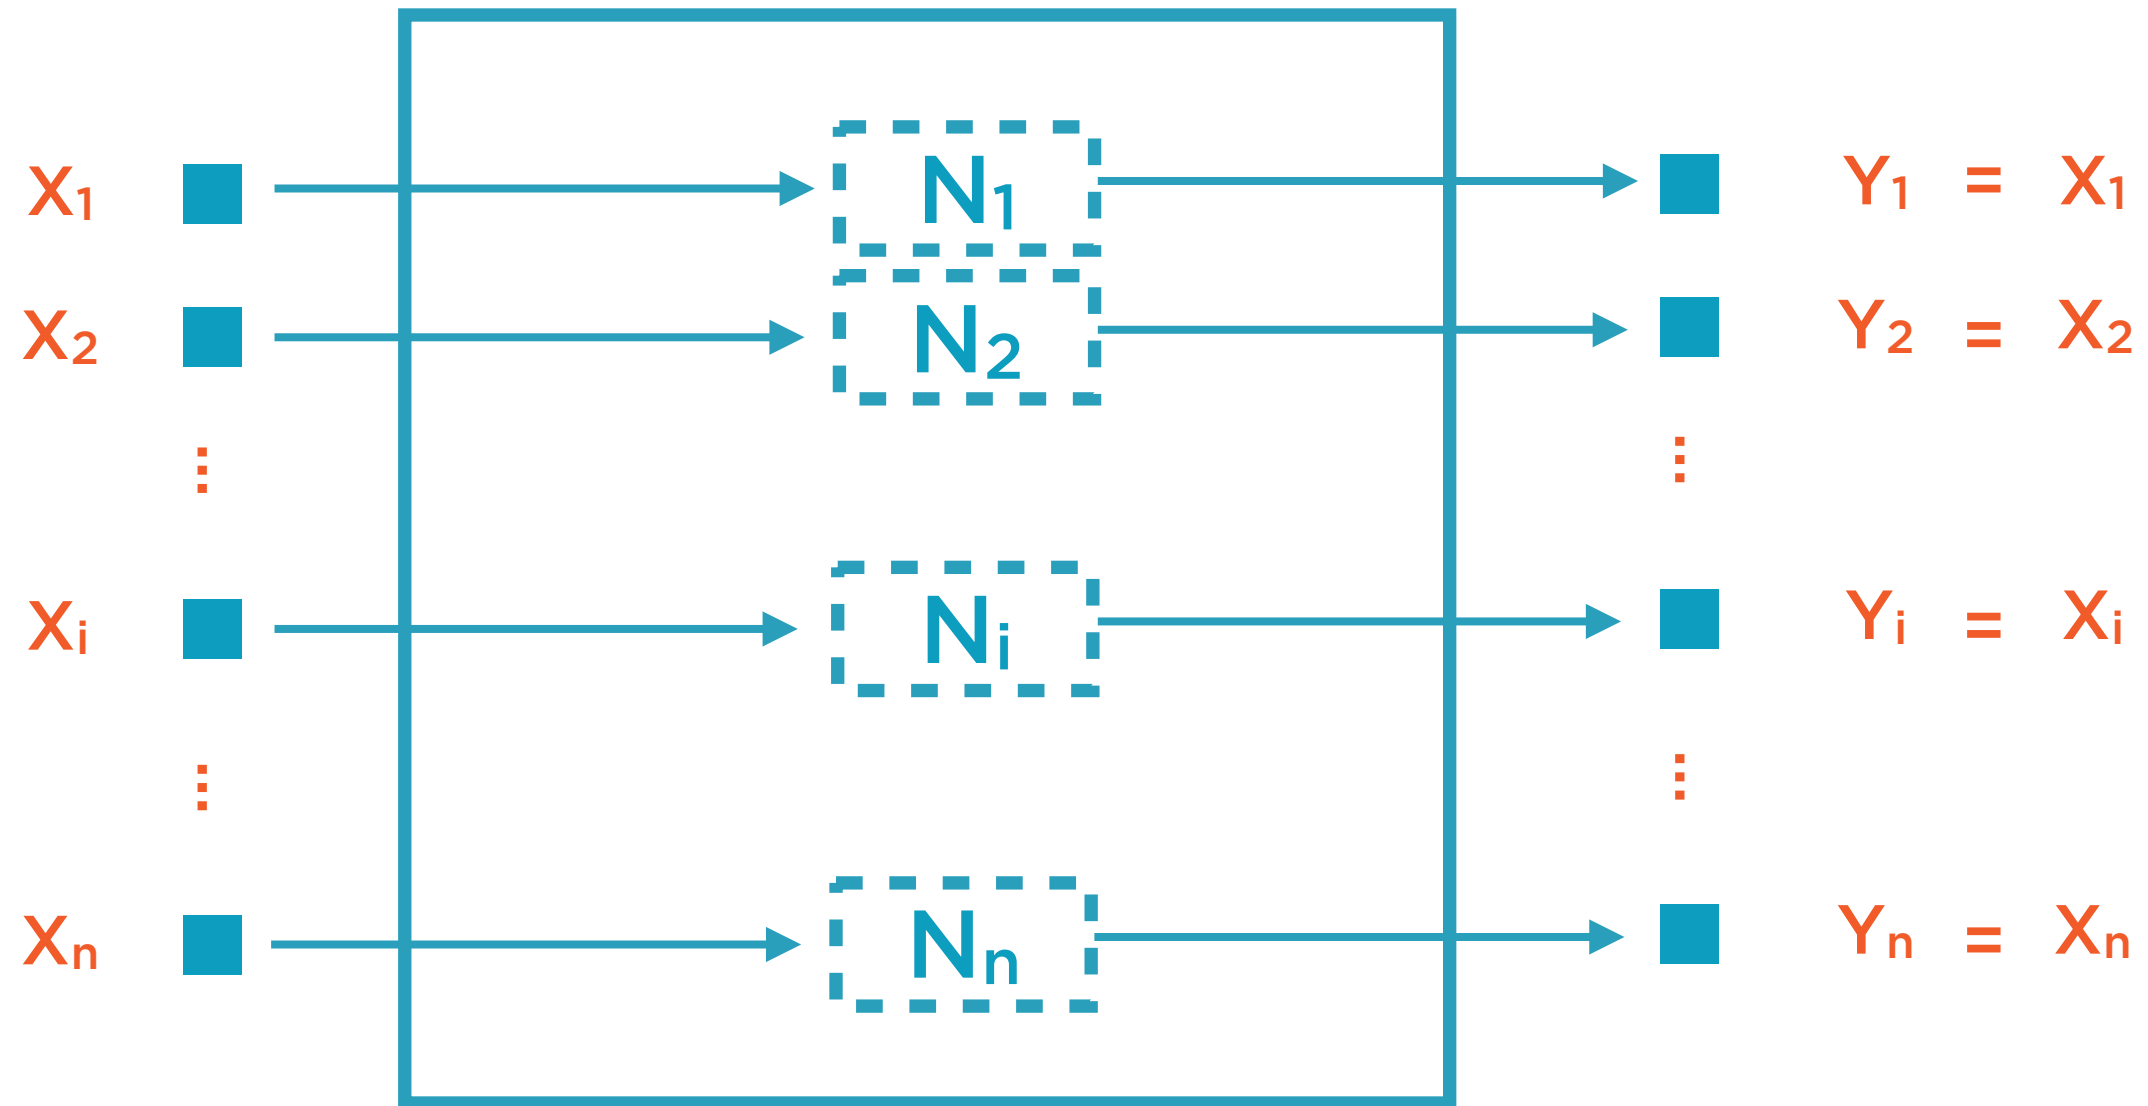
**Feedback (cost function) trains the model**

# Autoencoder



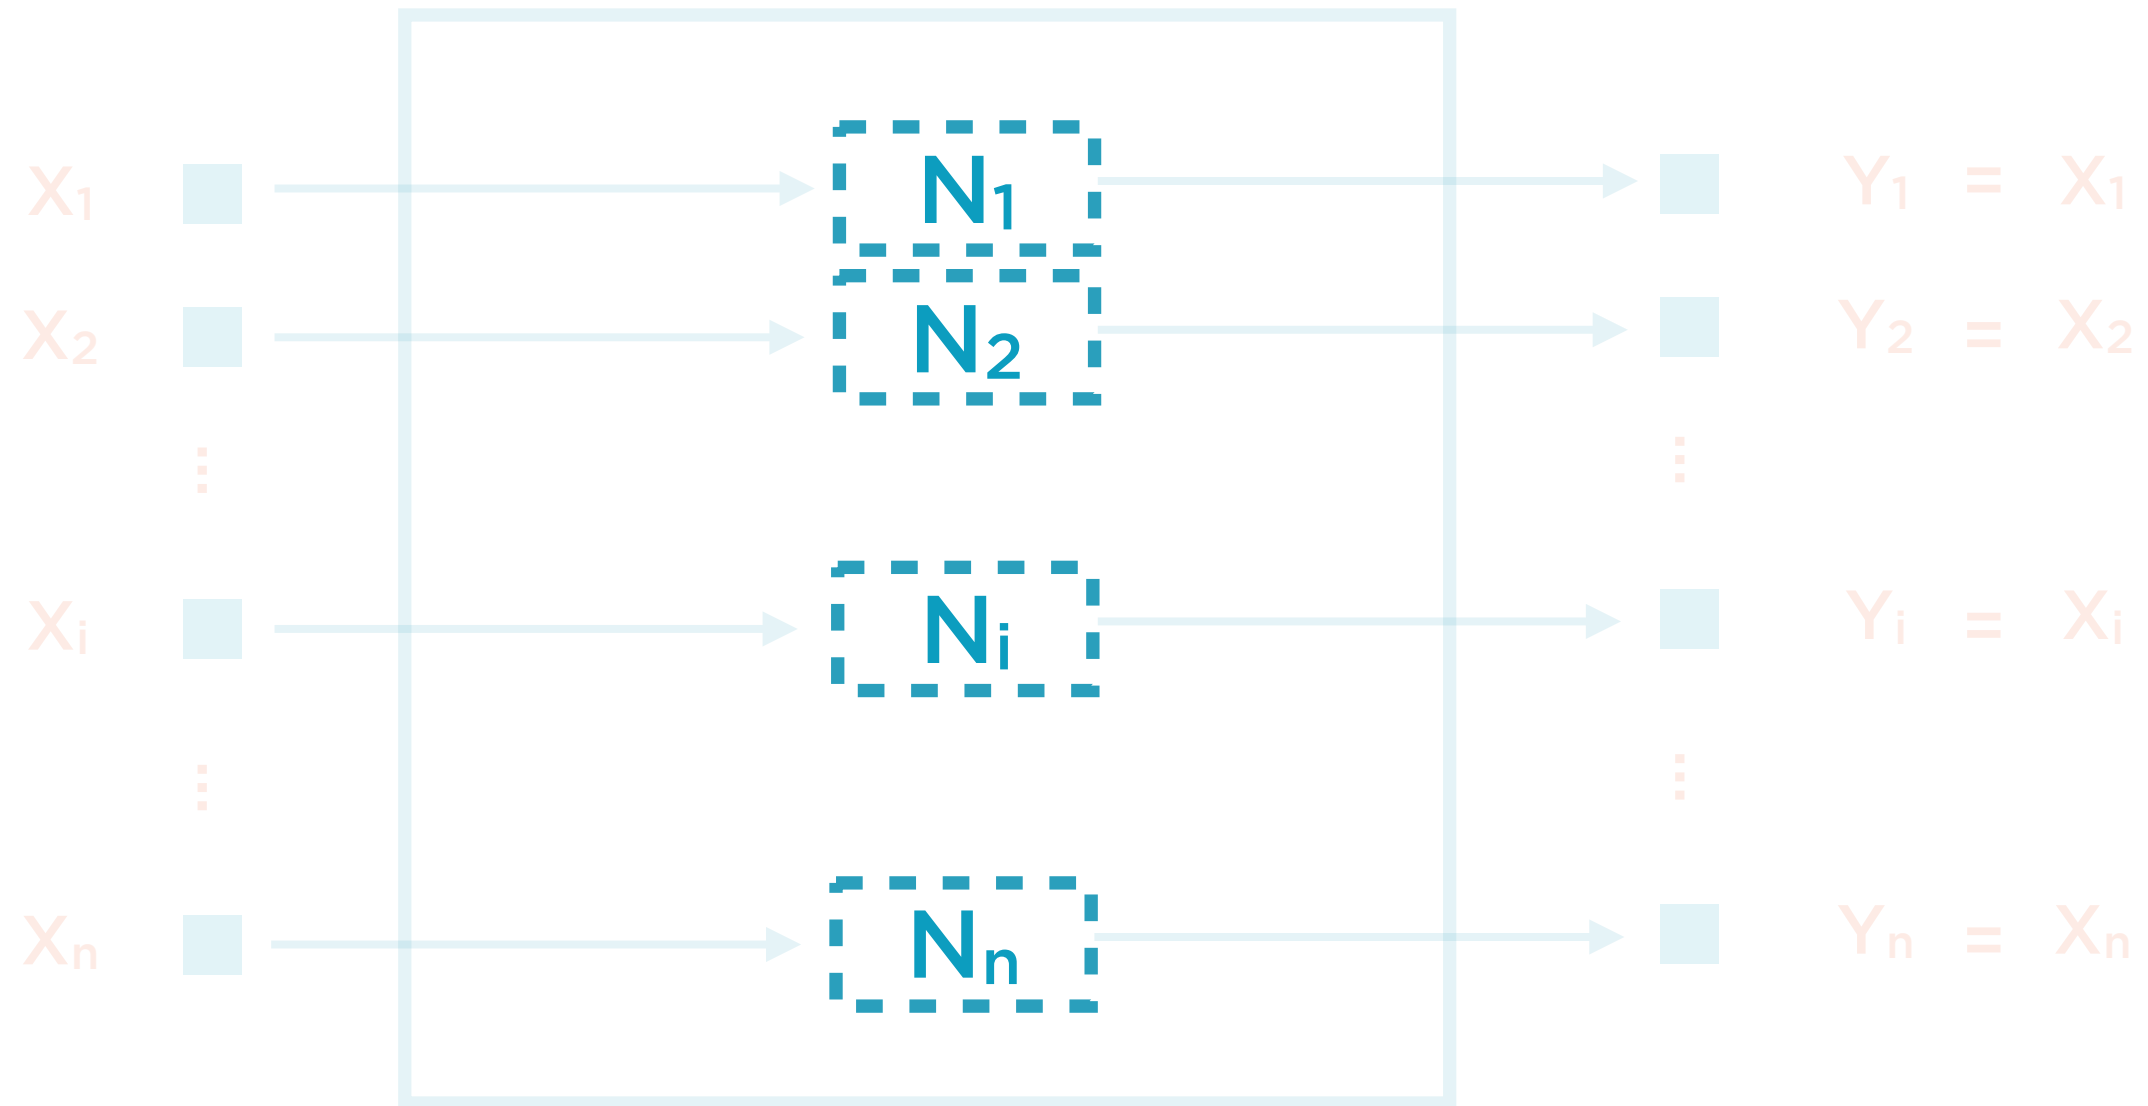
Autoencoders are Neural Networks  
that learn efficient representations  
of data (e.g. PCA)

# Trivial Autoencoder



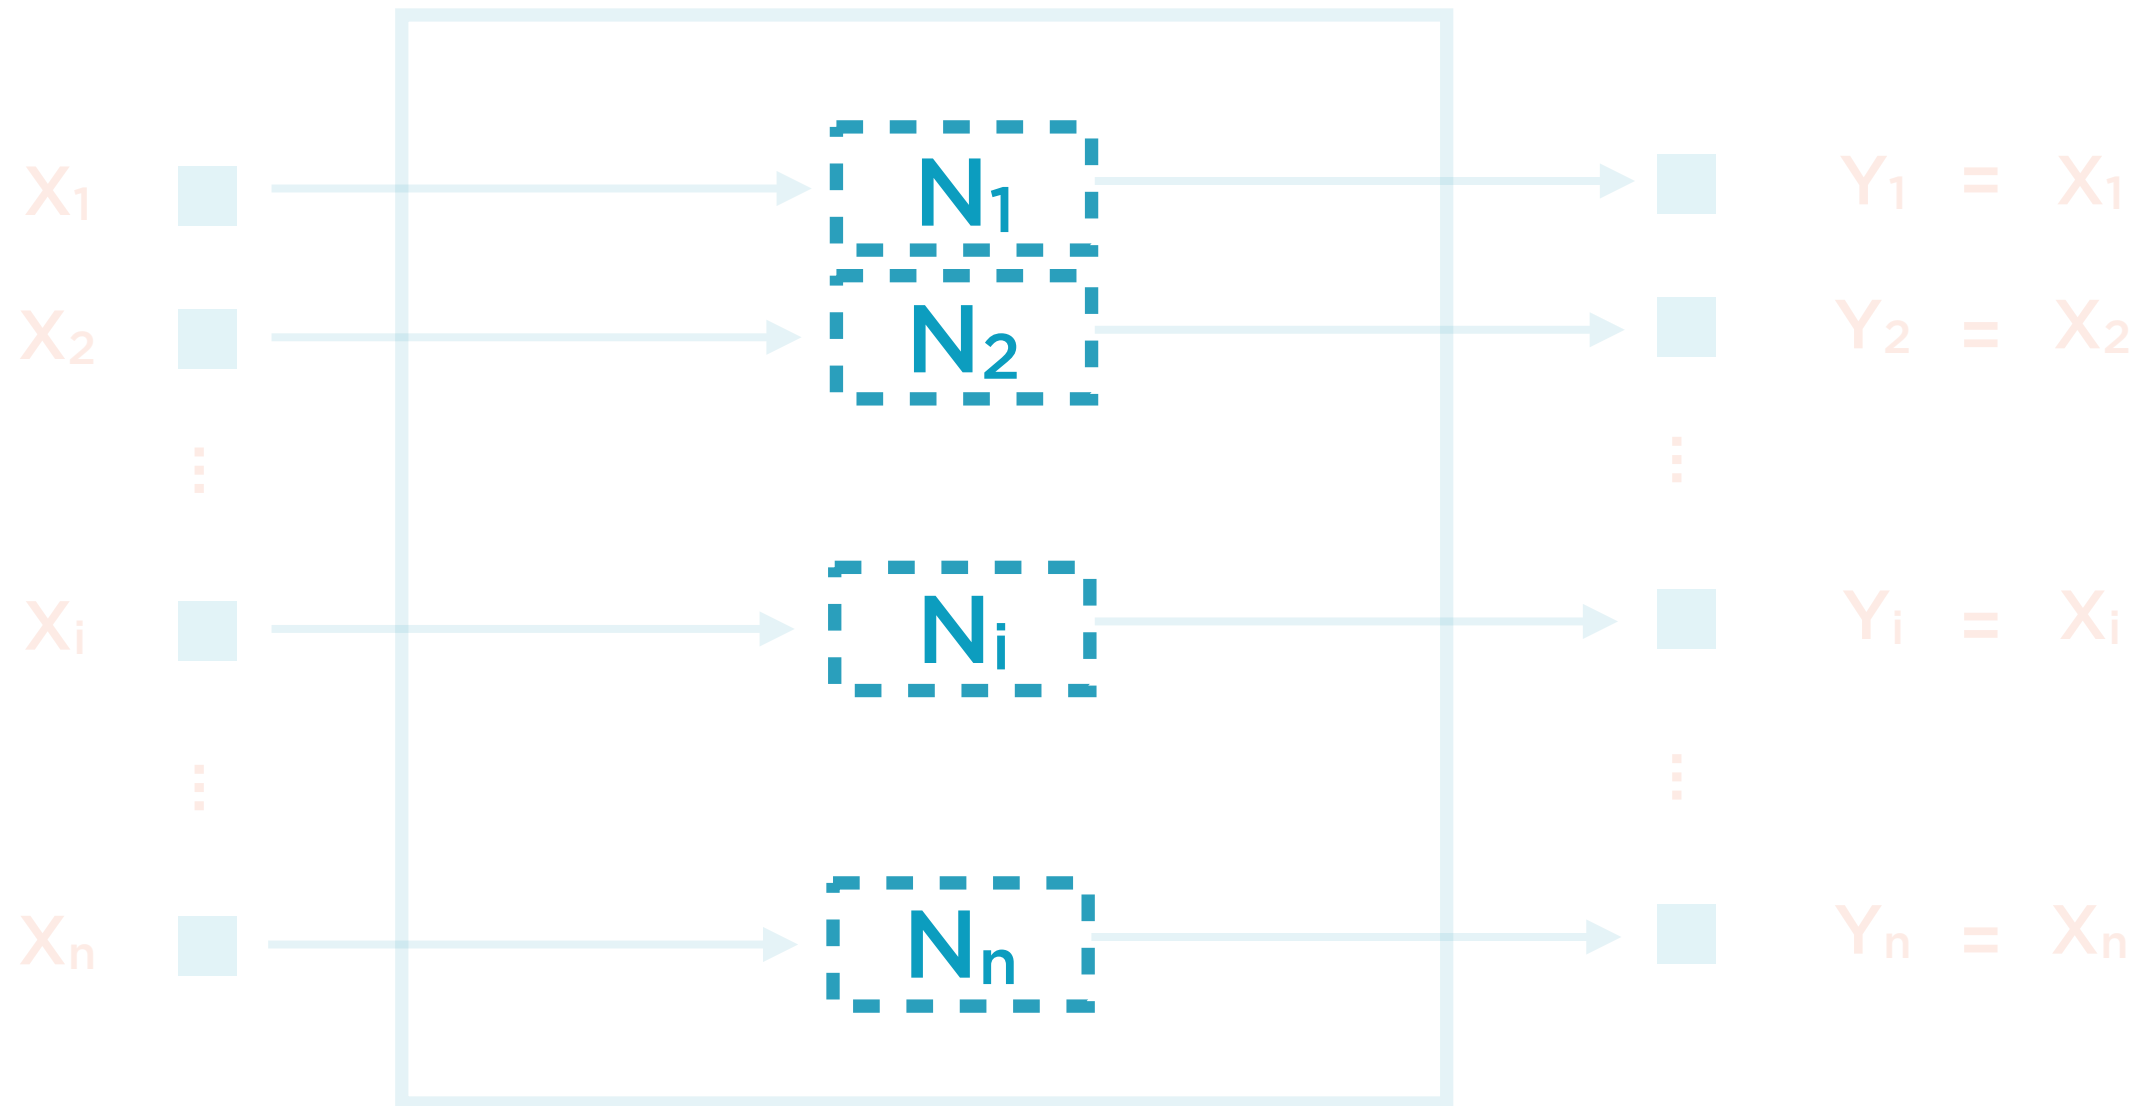
**This Neural Network trivially “learns” the input**

# Trivial Autoencoder



**Just one layer, which serves as both input and output layer**

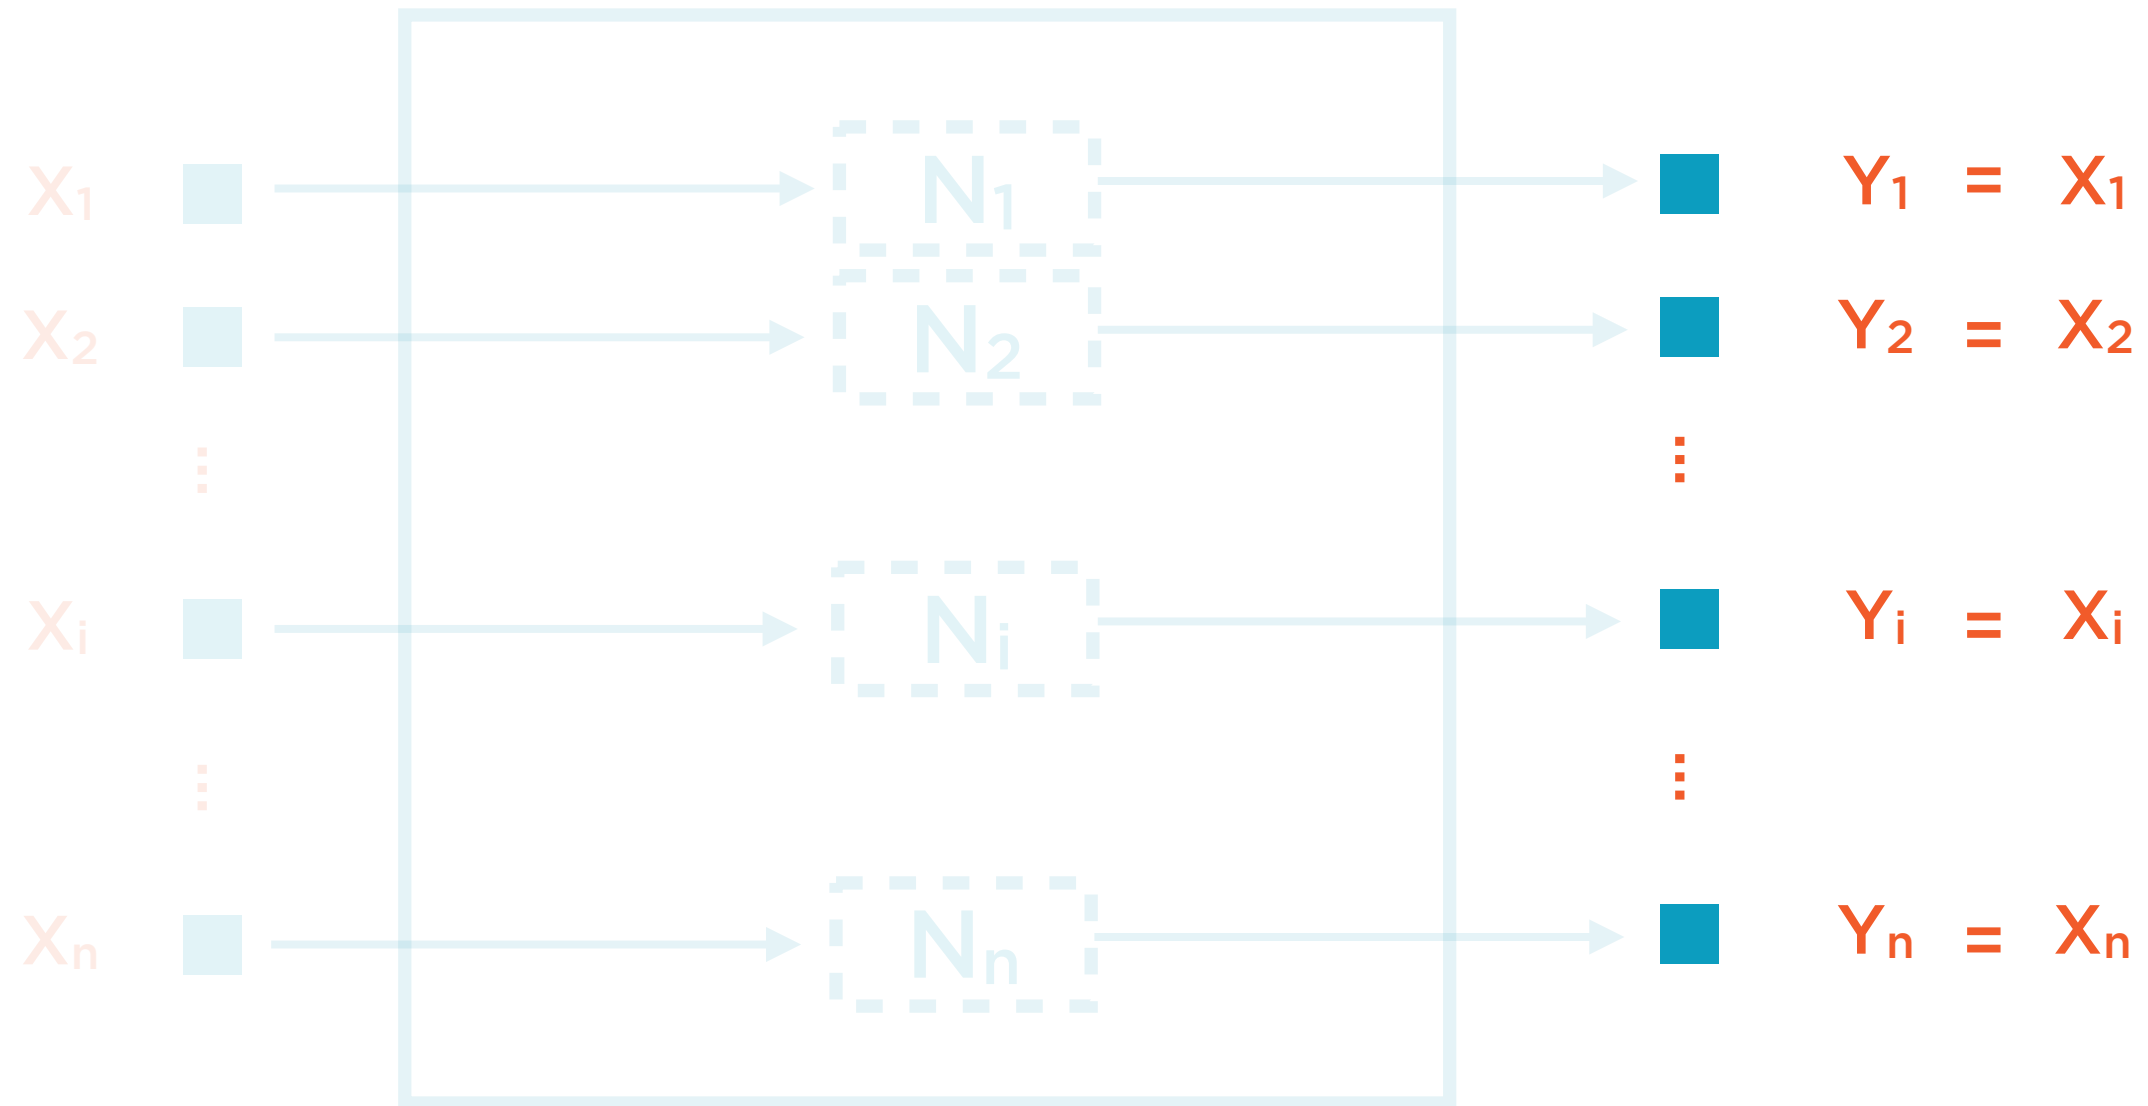
# Trivial Autoencoder



**In an autoencoder, the input and output layer must have same dimensionality as the input data**

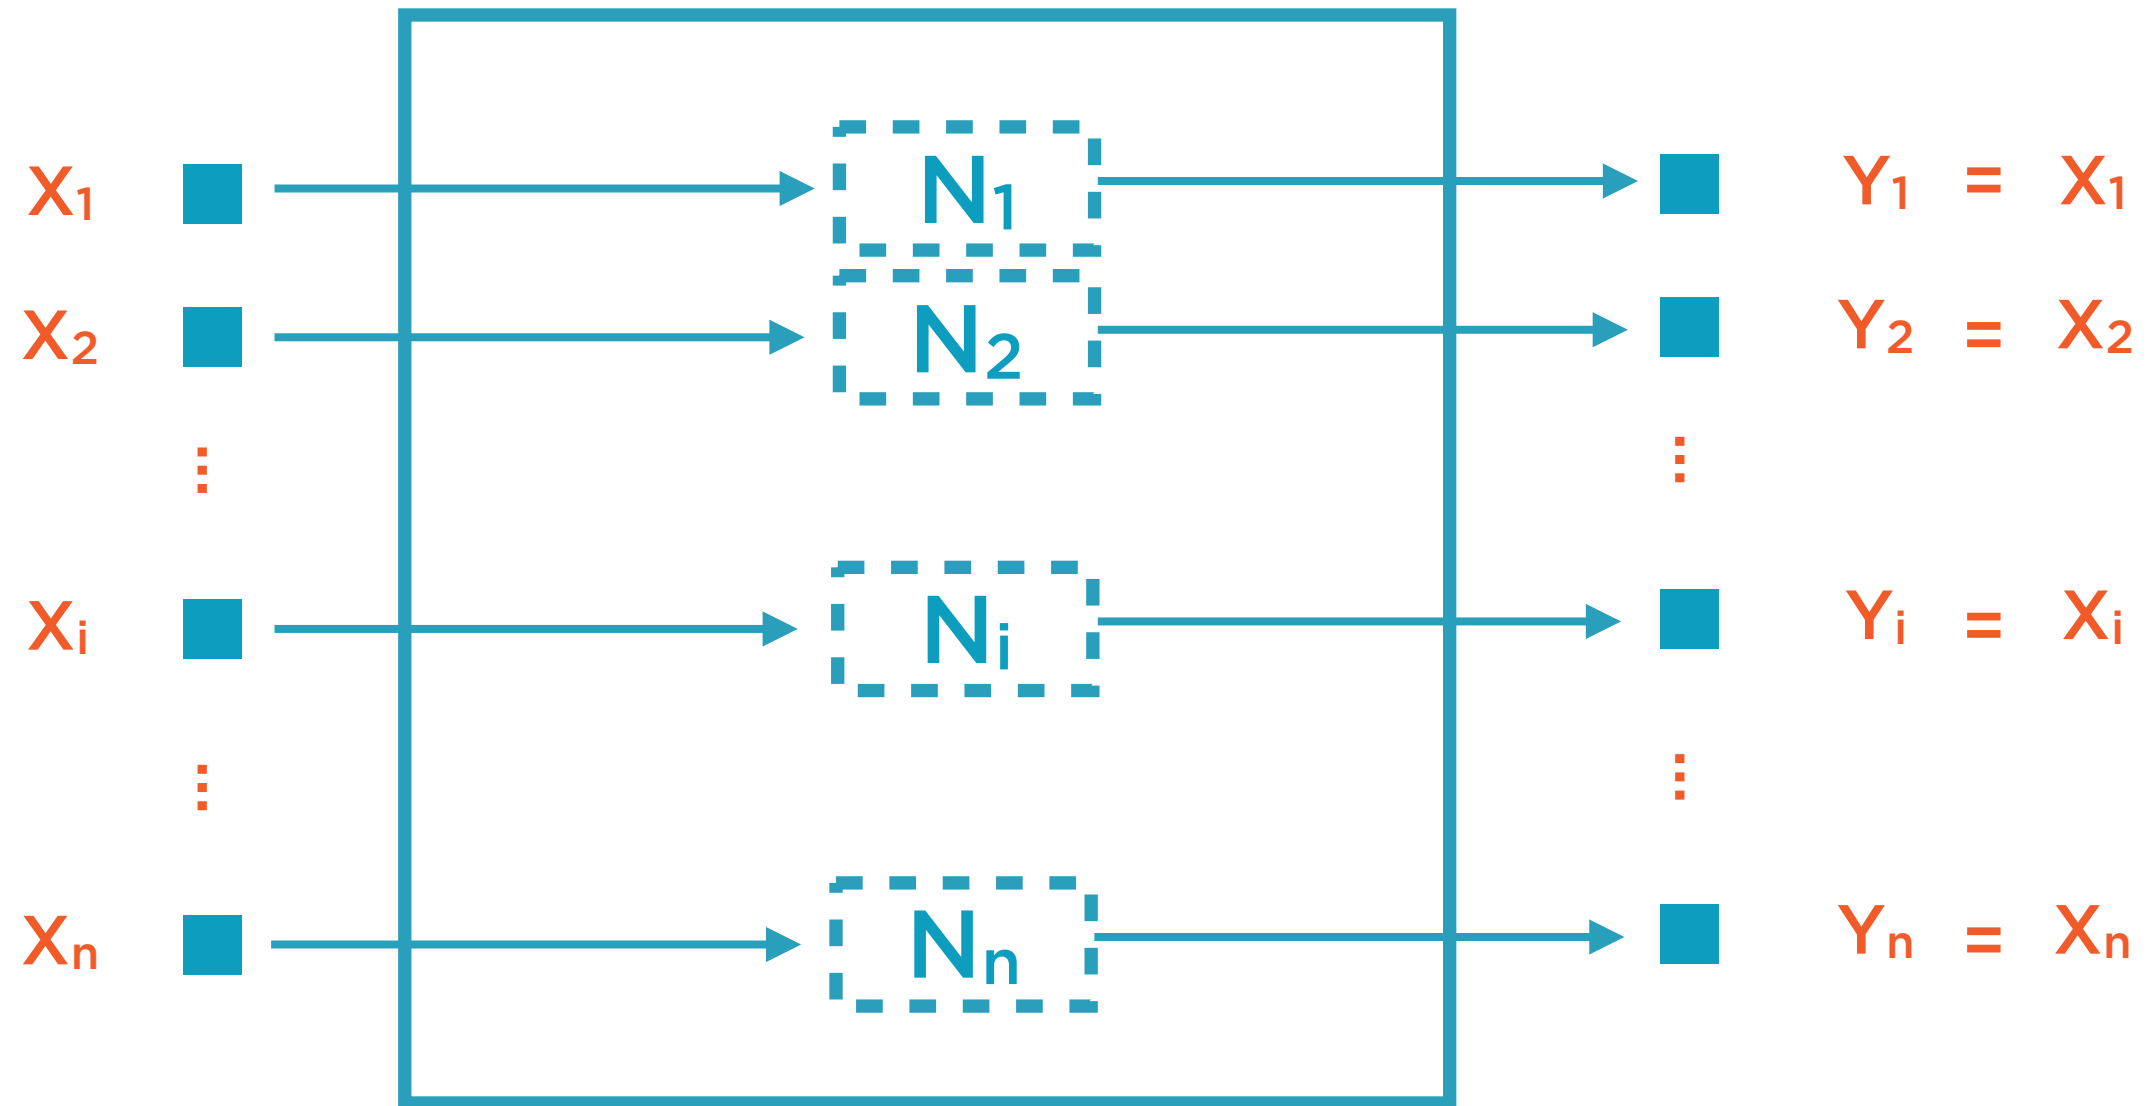


# Trivial Autoencoder



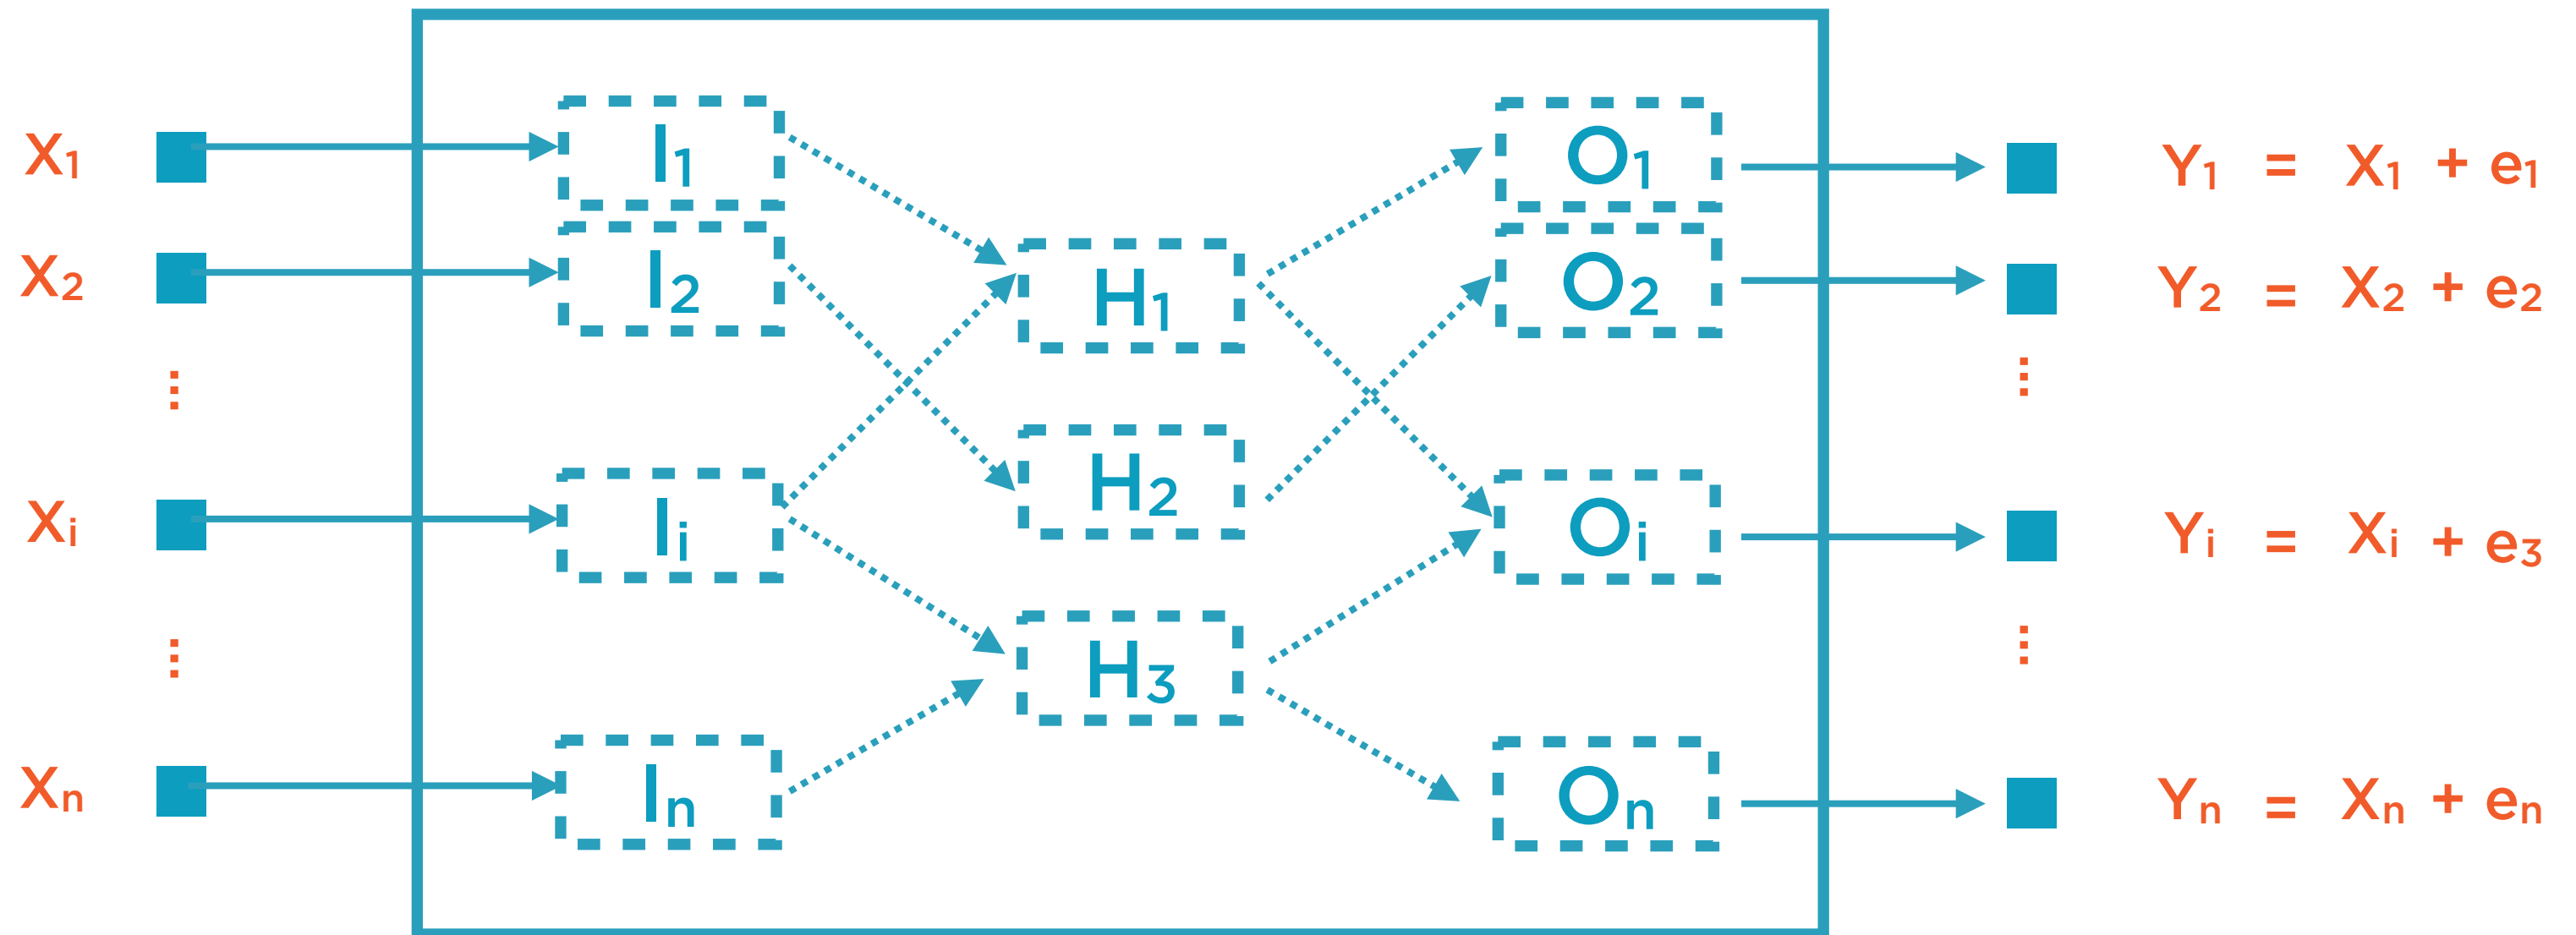
**The autoencoder just passes the input through, so output is exactly equal to input**

# Trivial Autoencoder



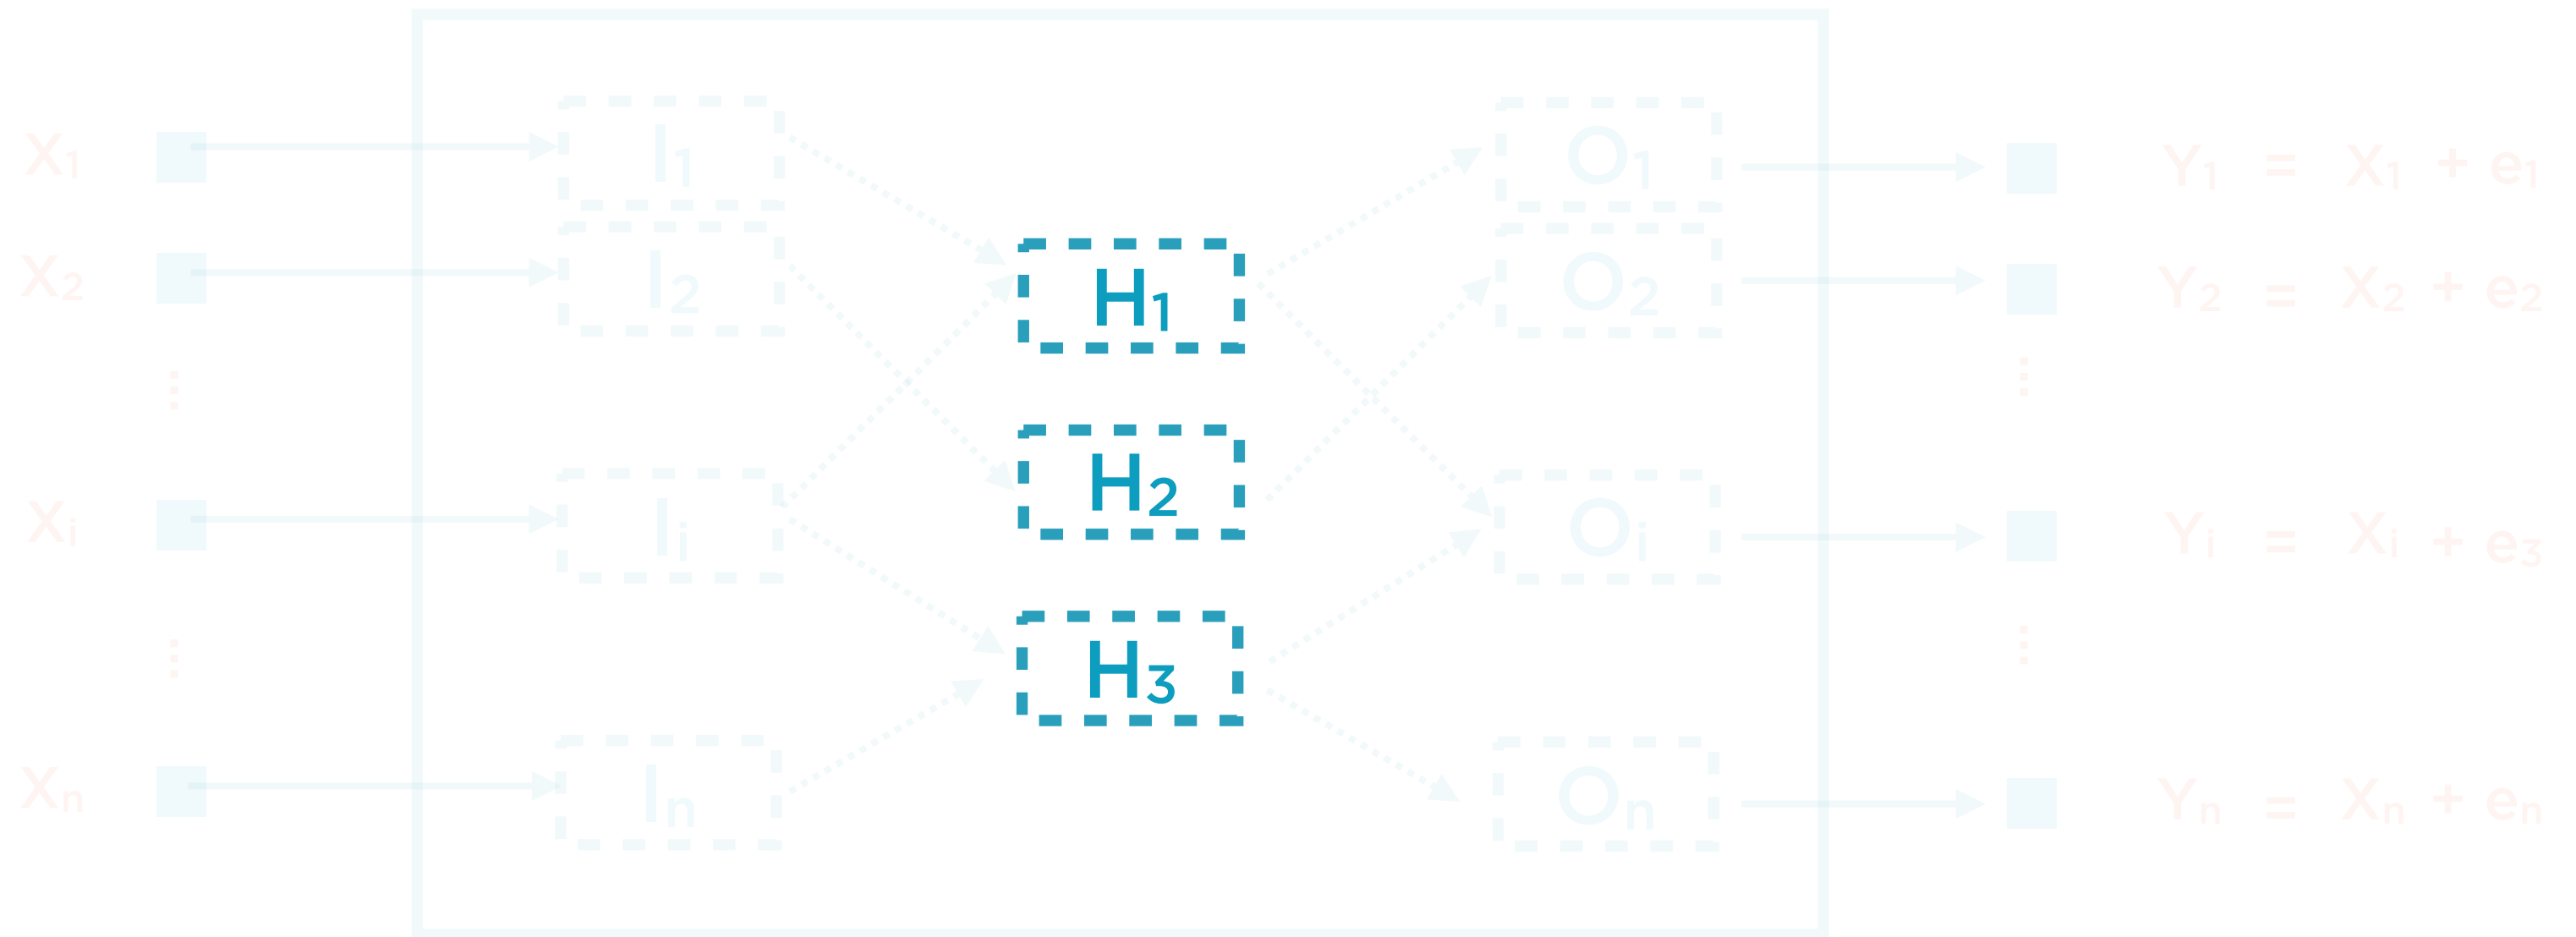
**Now, let's constrain the network to force dimensionality reduction**

# Undercomplete Autoencoder



**Dimensionality of the NN is now lower than that of input data**

# Undercomplete Autoencoder



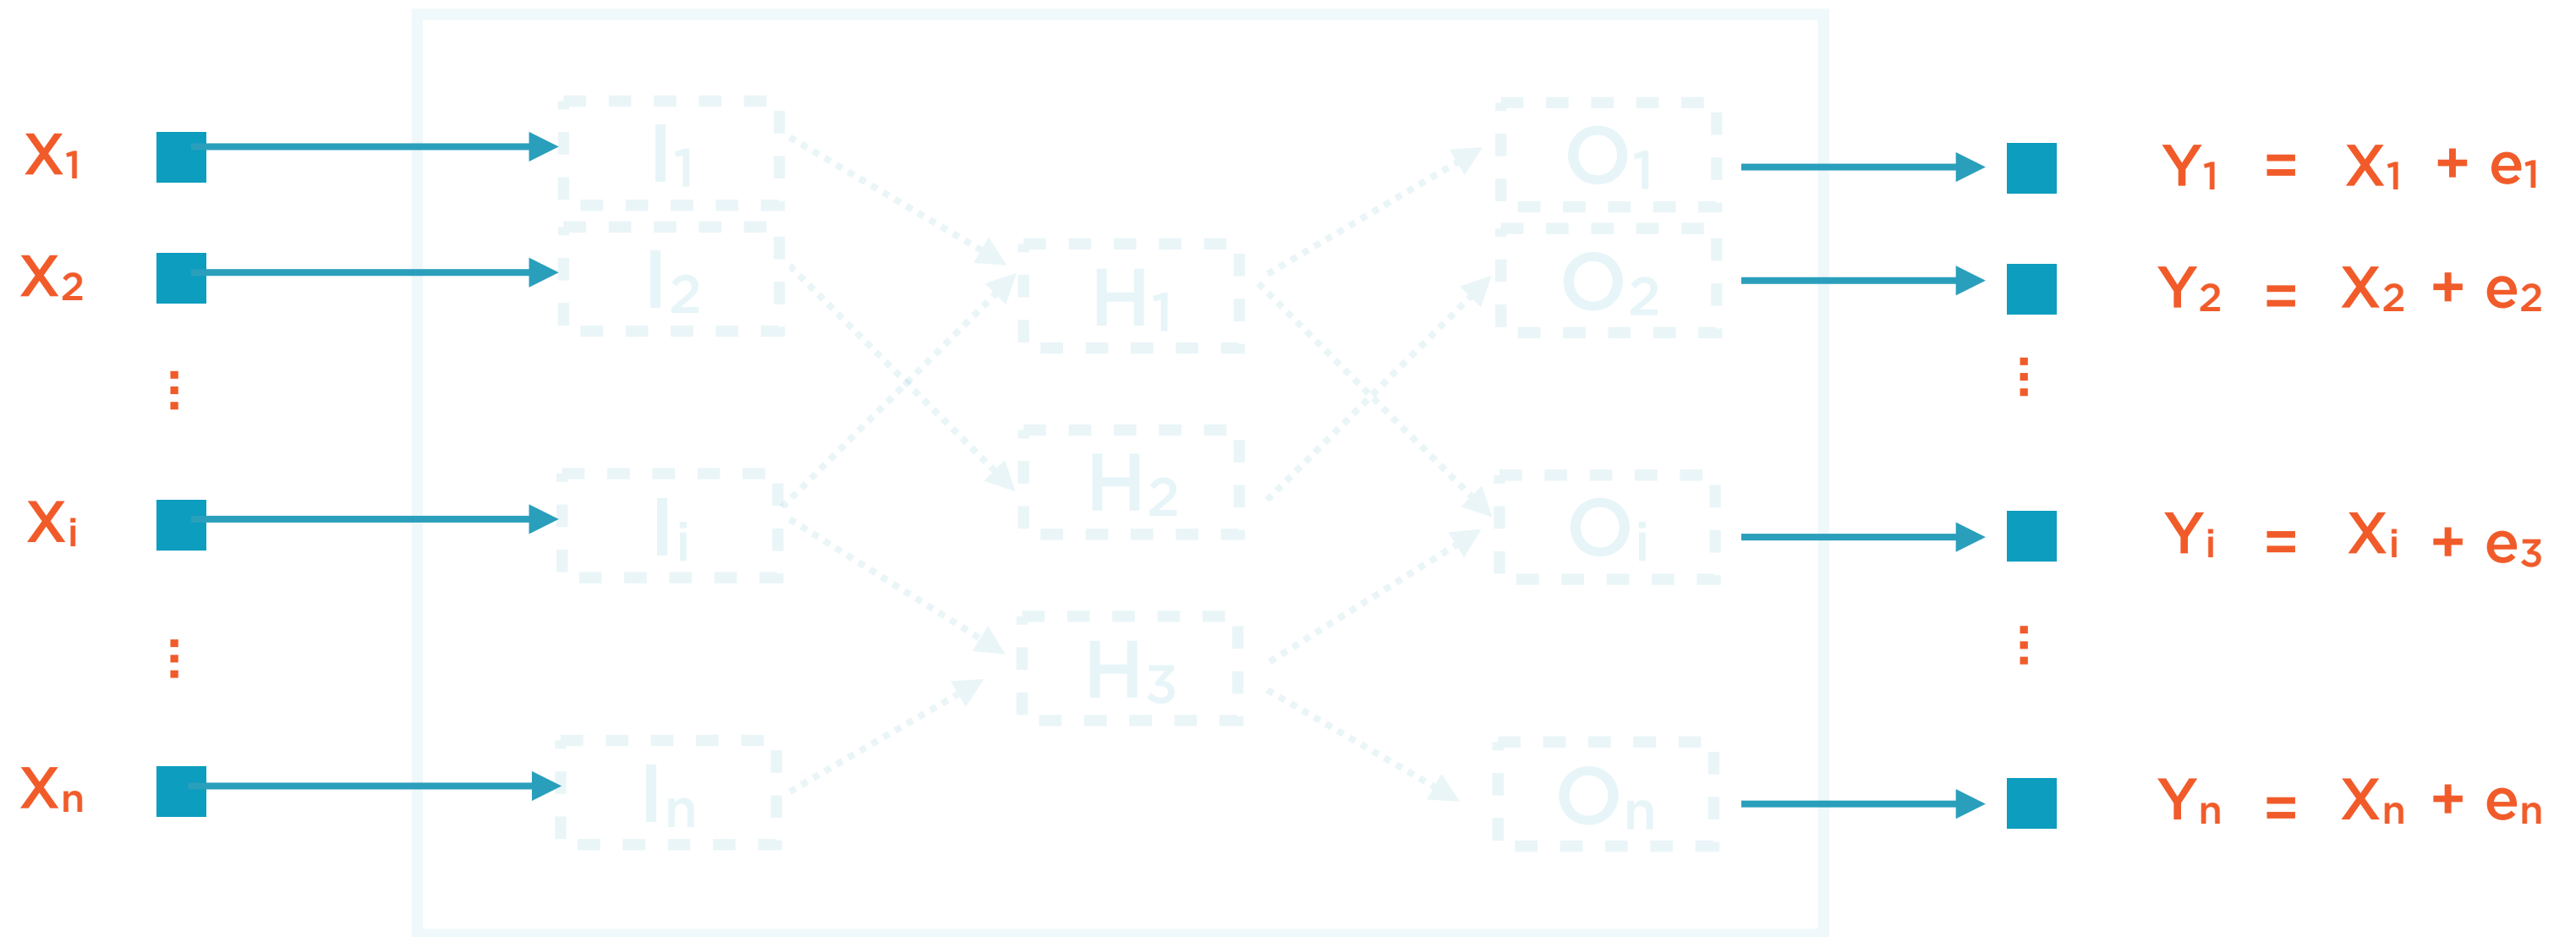
**Add a middle, hidden layer with just three neurons ( $3 < N$ )**

# Undercomplete Autoencoder



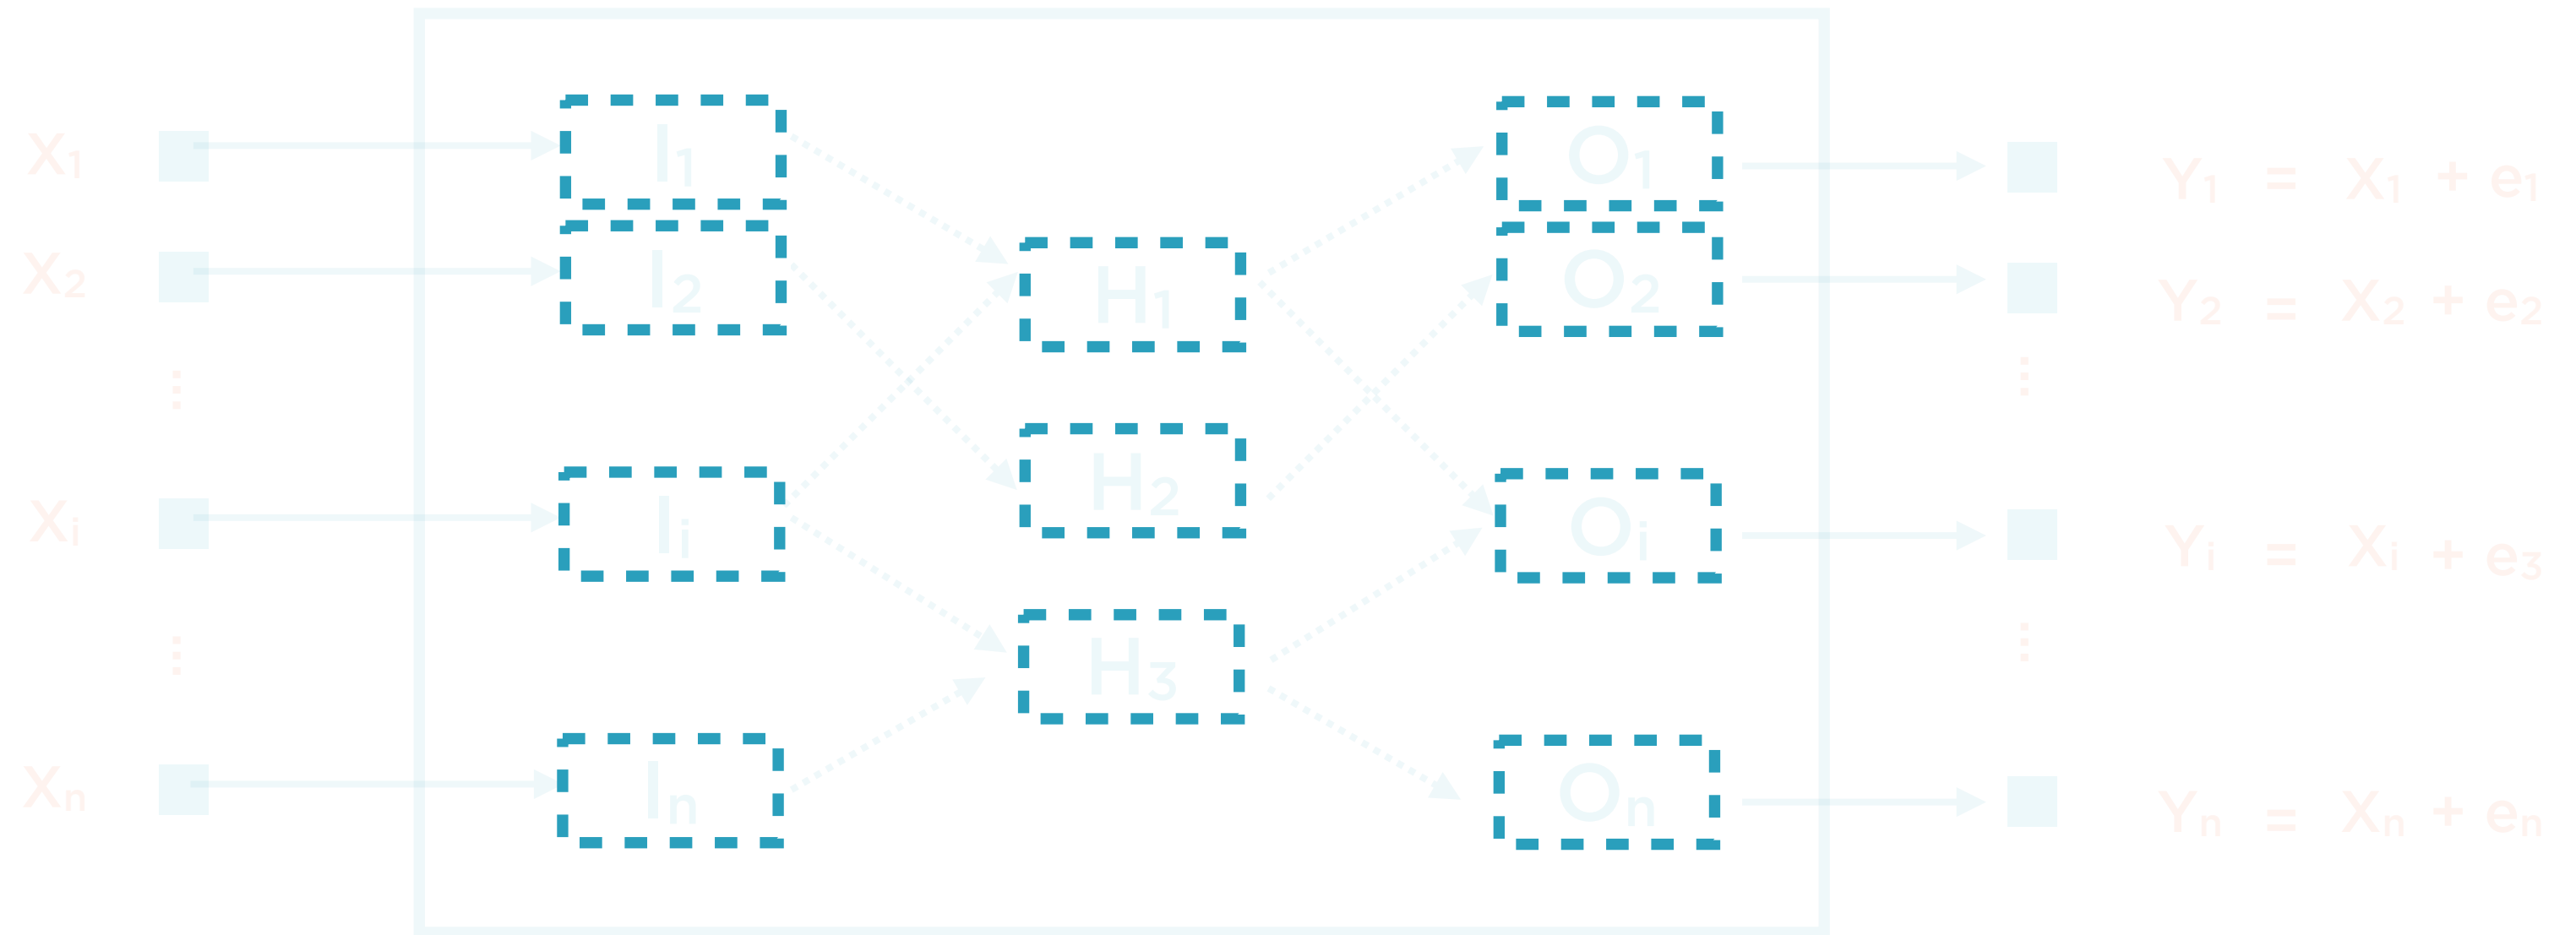
**The input and output layers must now be separated, since each must still have same dimensionality as the input**

# Undercomplete Autoencoder



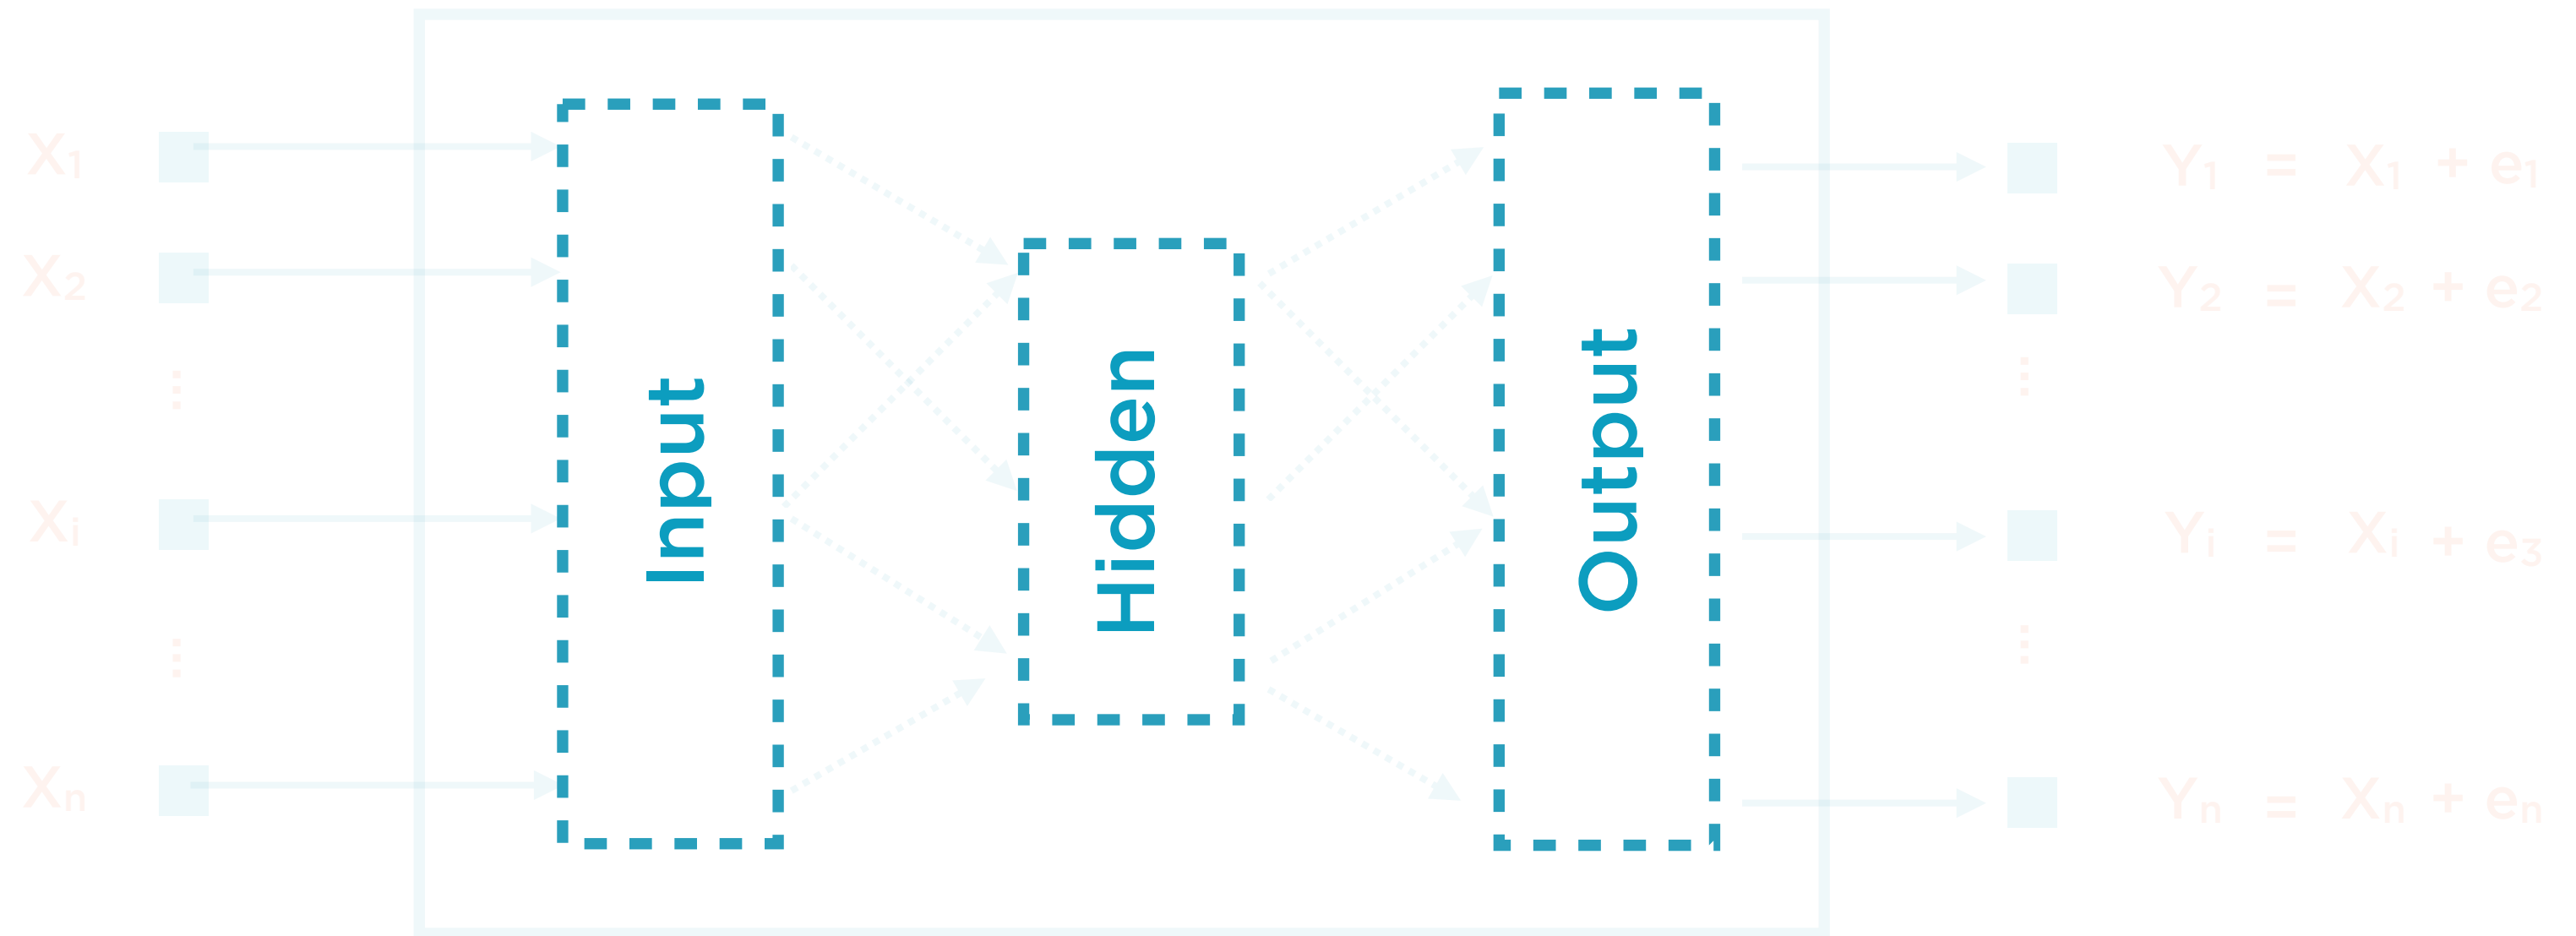
**Why? Because autoencoder seeks to reconstruct input**

# Undercomplete Autoencoder



**This gives undercomplete autoencoders a characteristic sandwich-like appearance**

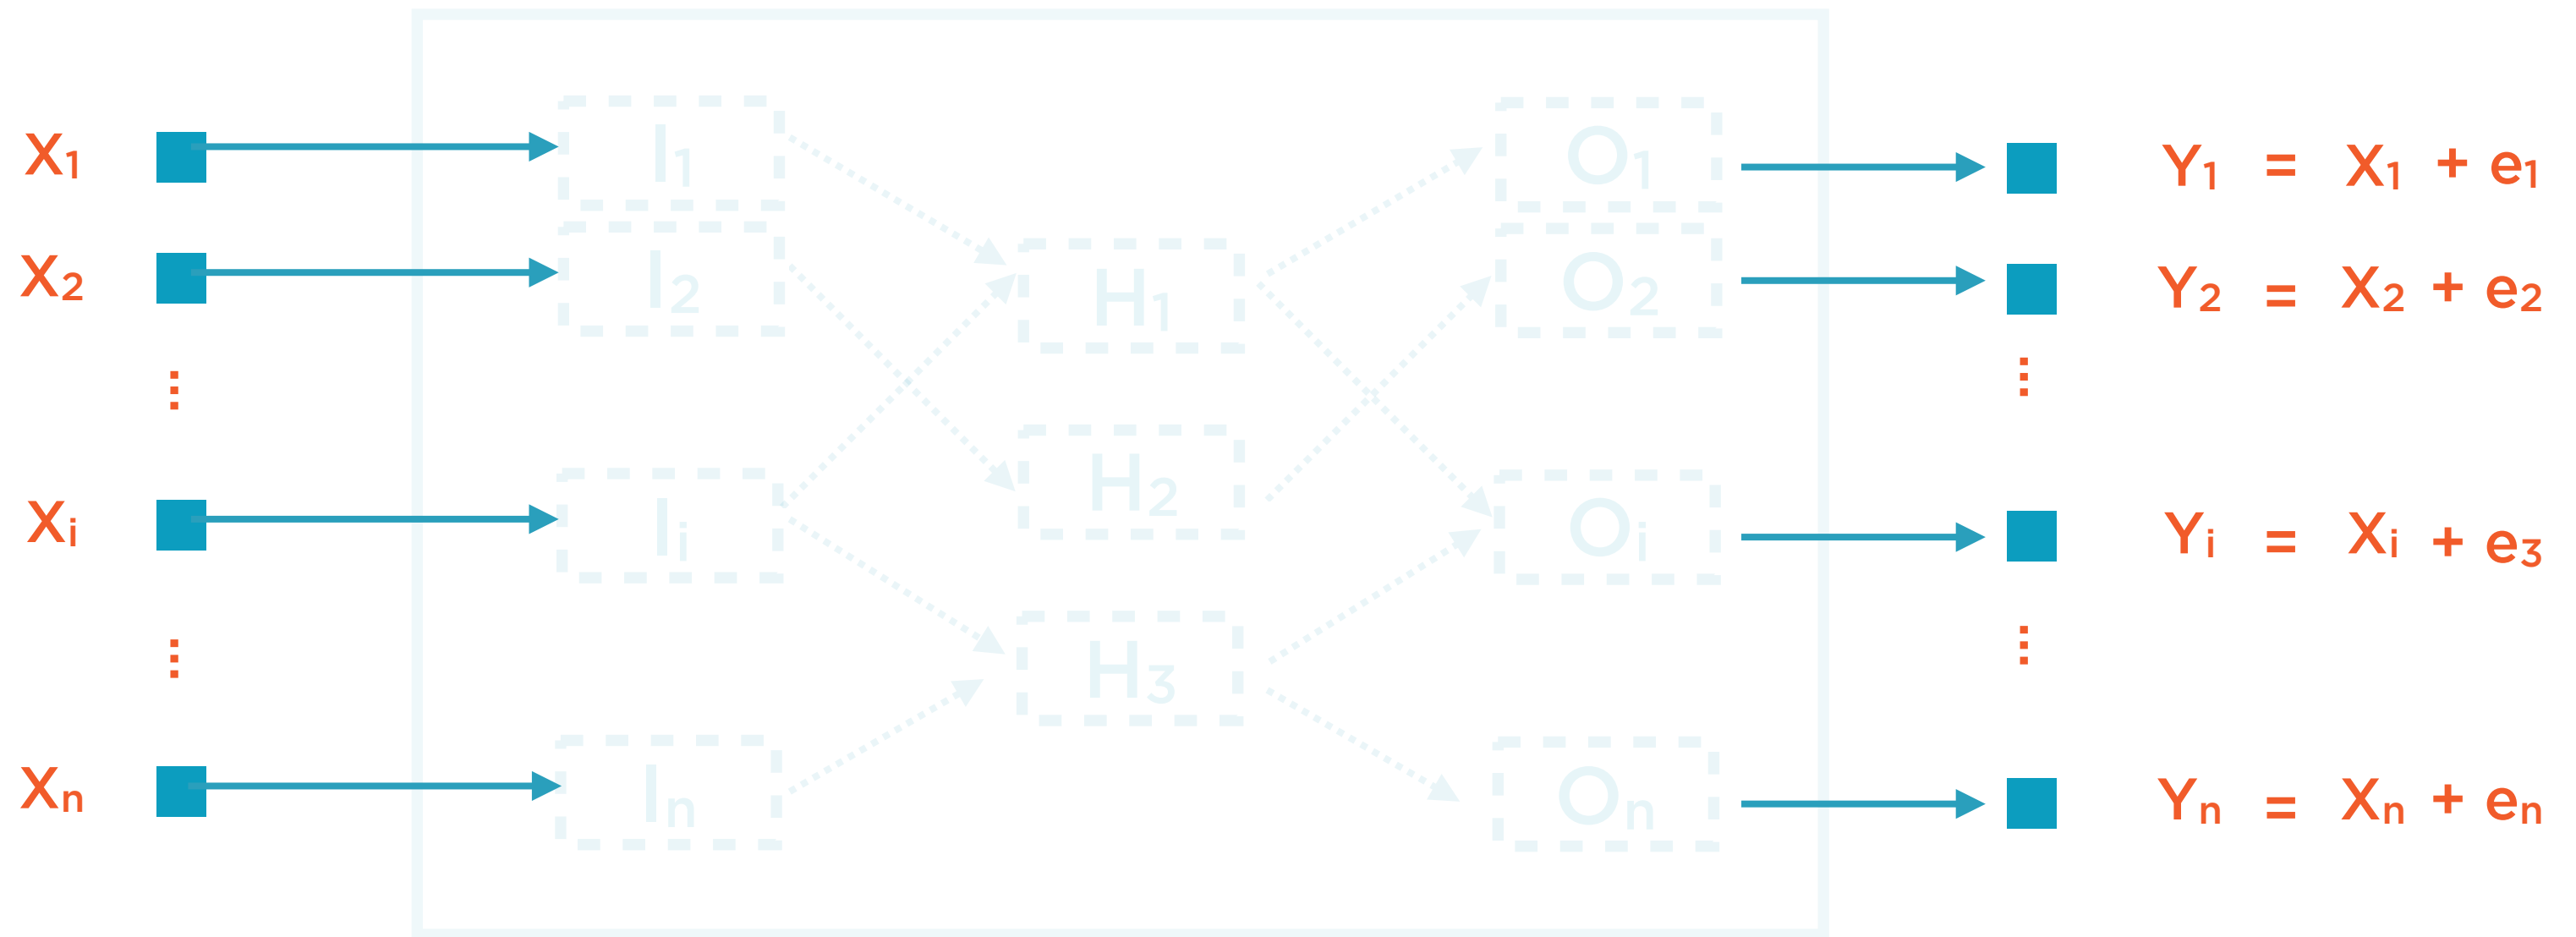
# Undercomplete Autoencoder



**This gives undercomplete autoencoders a characteristic sandwich-like appearance**

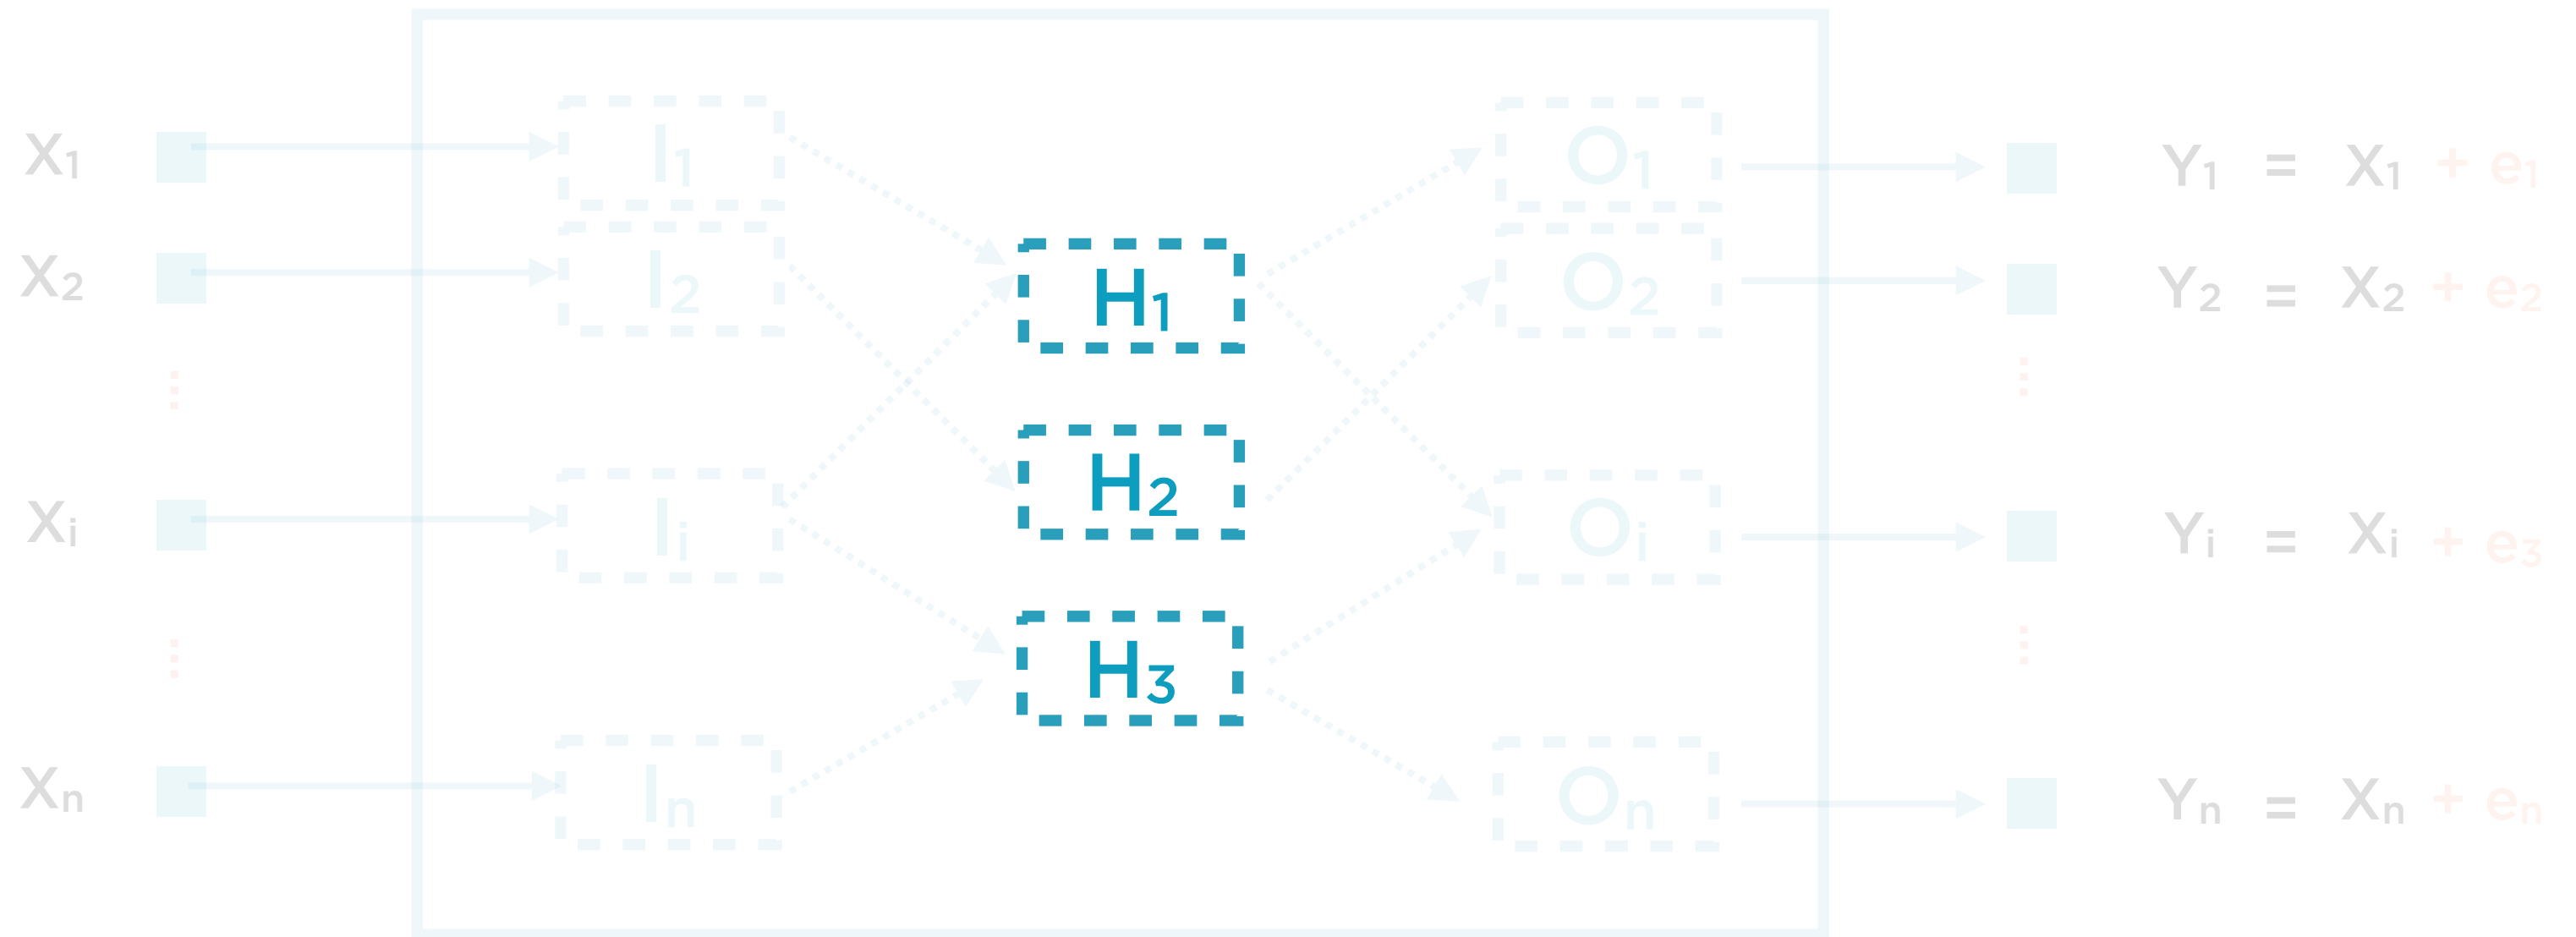


# Undercomplete Autoencoder



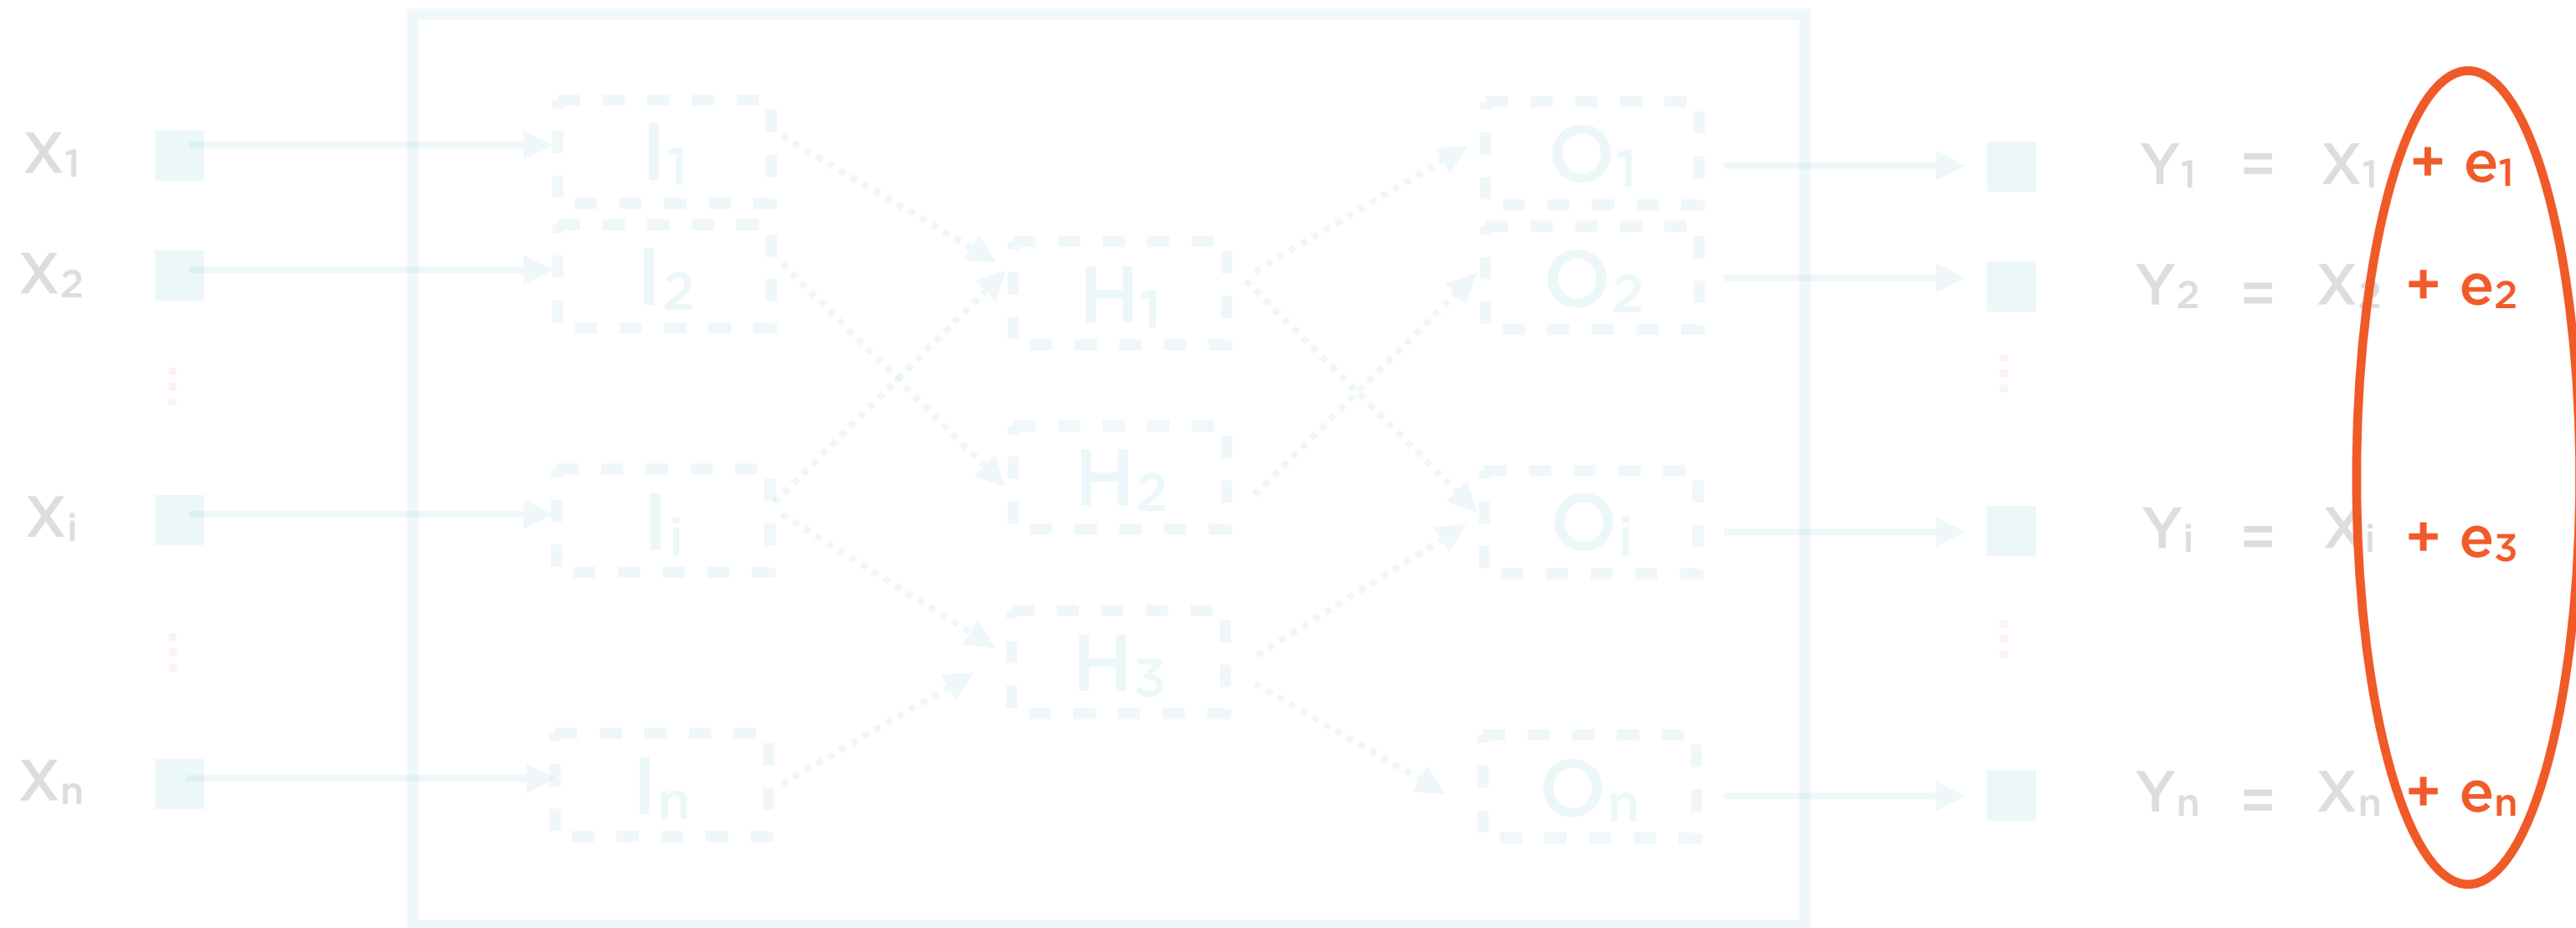
**The undercomplete autoencoder will try to exactly match the input,  
but it will likely not succeed completely**

# Undercomplete Autoencoder



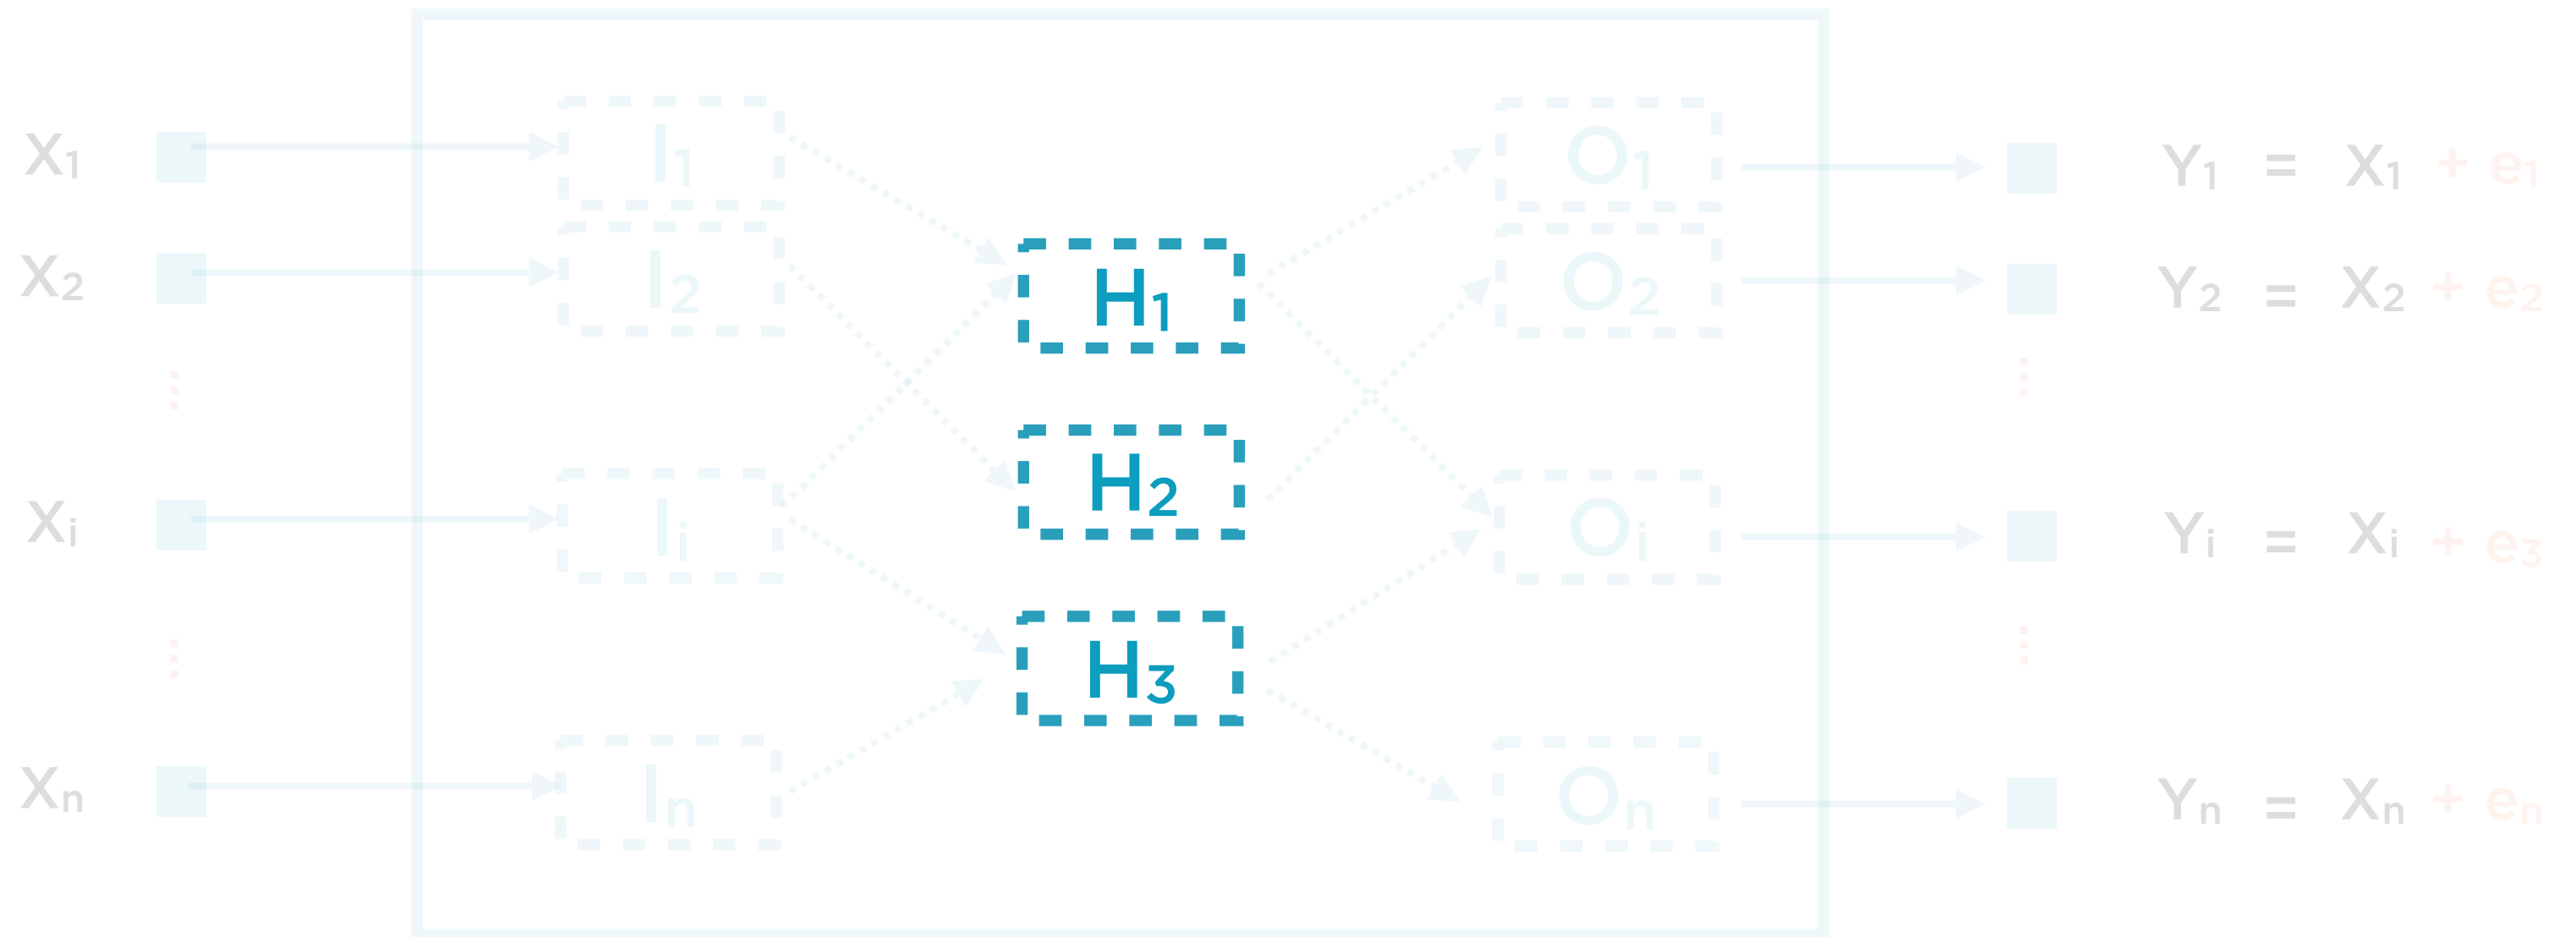
Now, because of the dimensionality reduction, output will **not** be exactly same as input

# Undercomplete Autoencoder



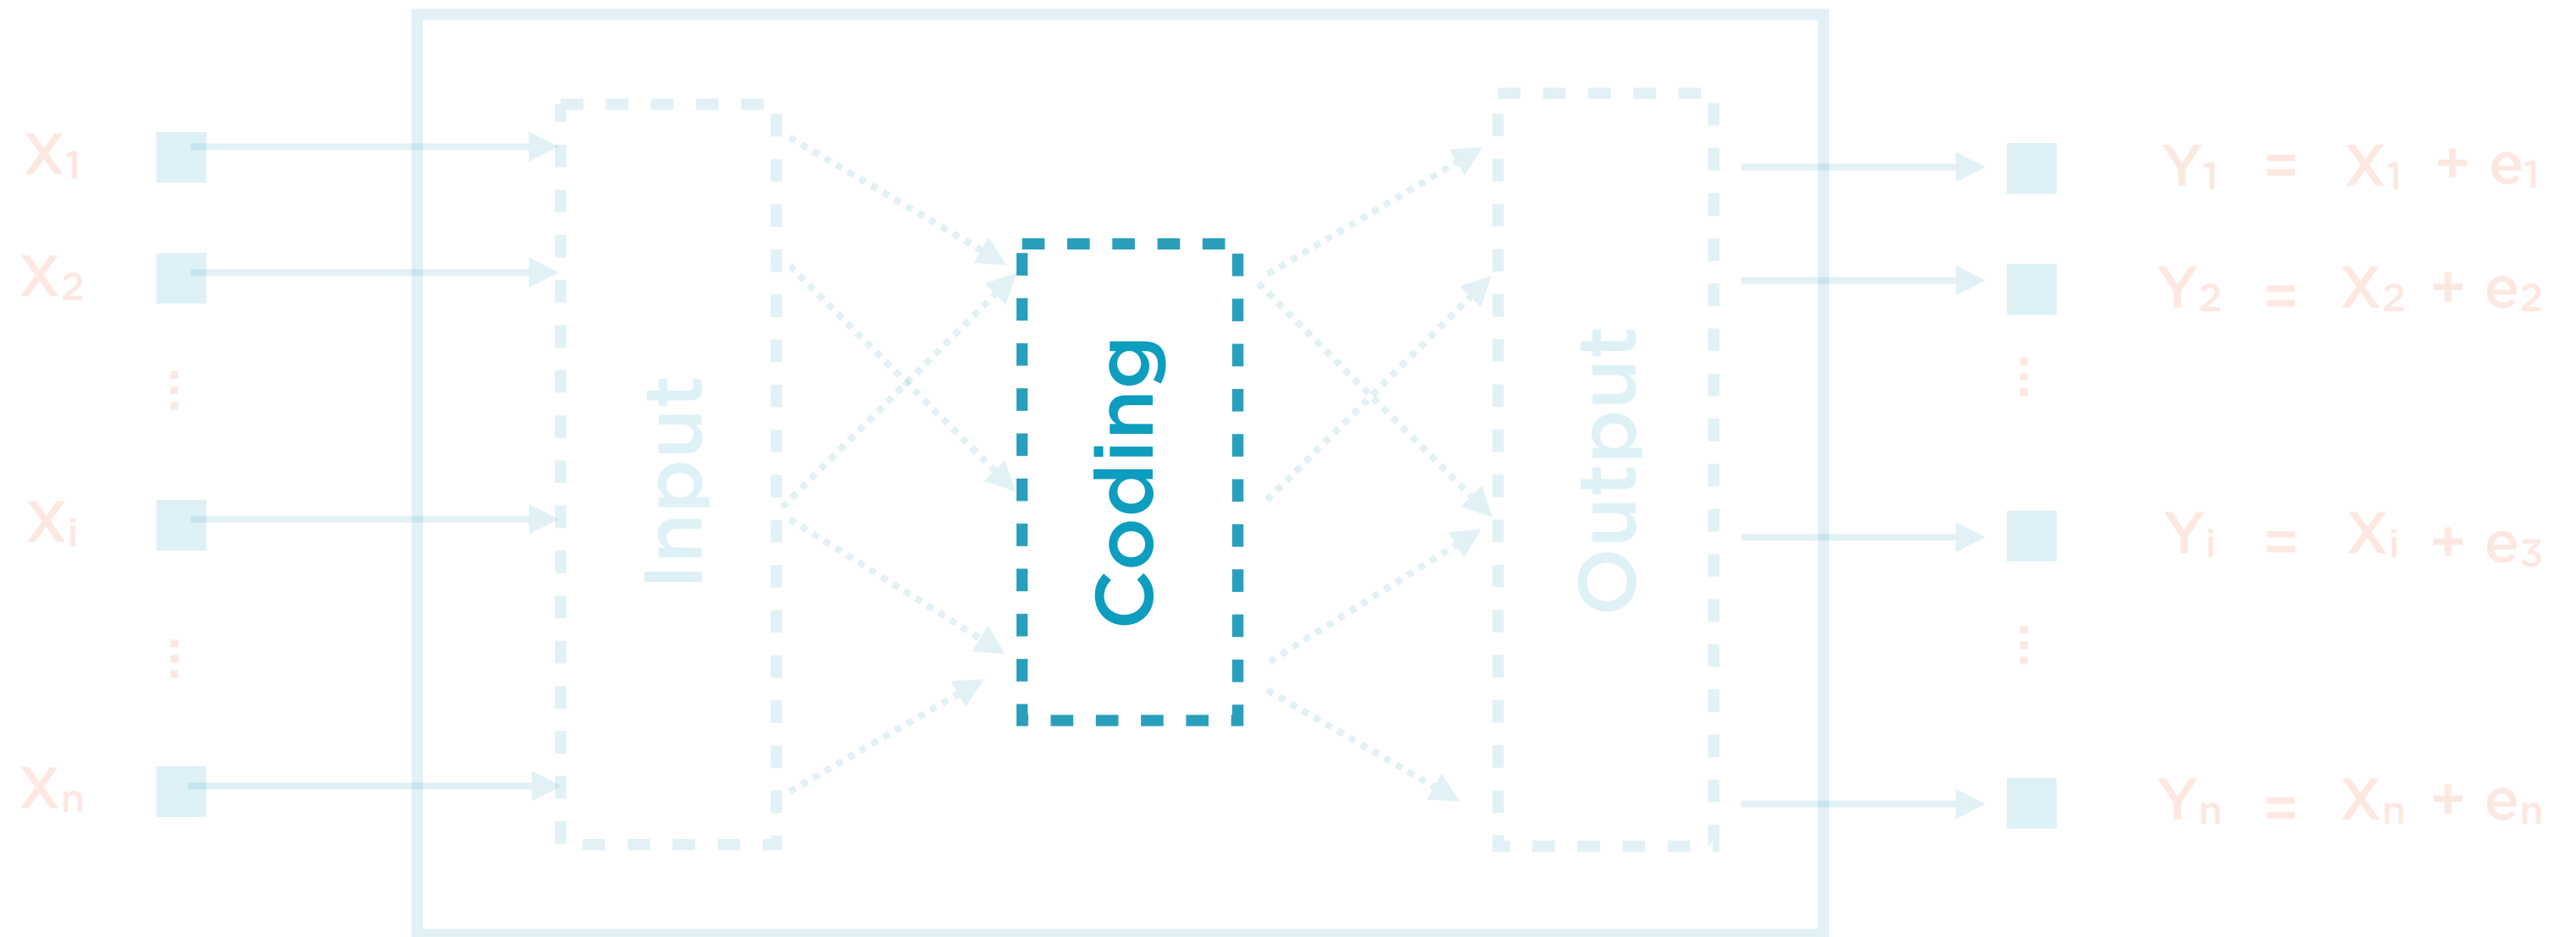
**A reconstruction error will now exist**

# Undercomplete Autoencoder



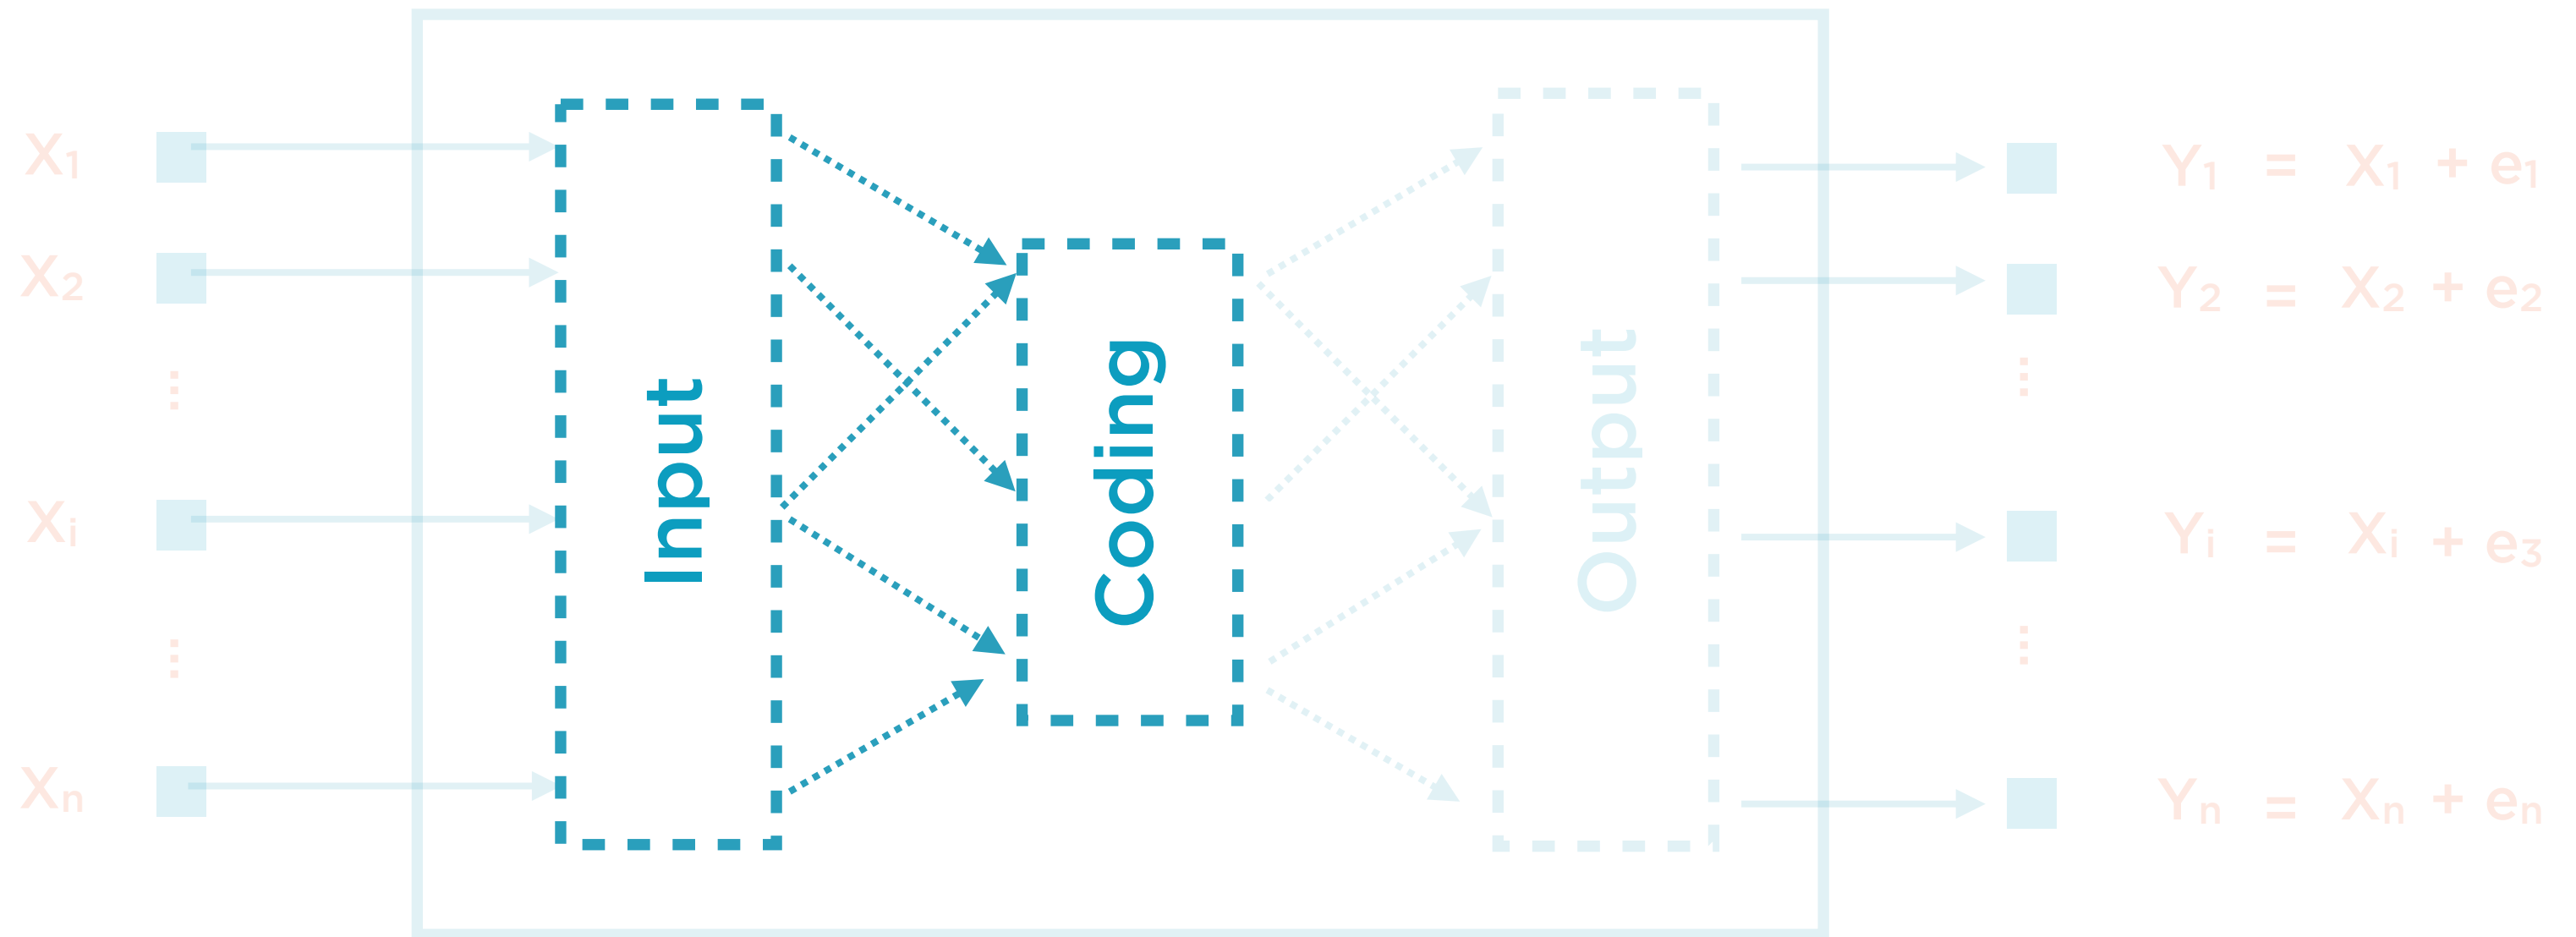
The autoencoder will be forced to **learn the most significant characteristics** of the data i.e. latent factors

# Undercomplete Autoencoder



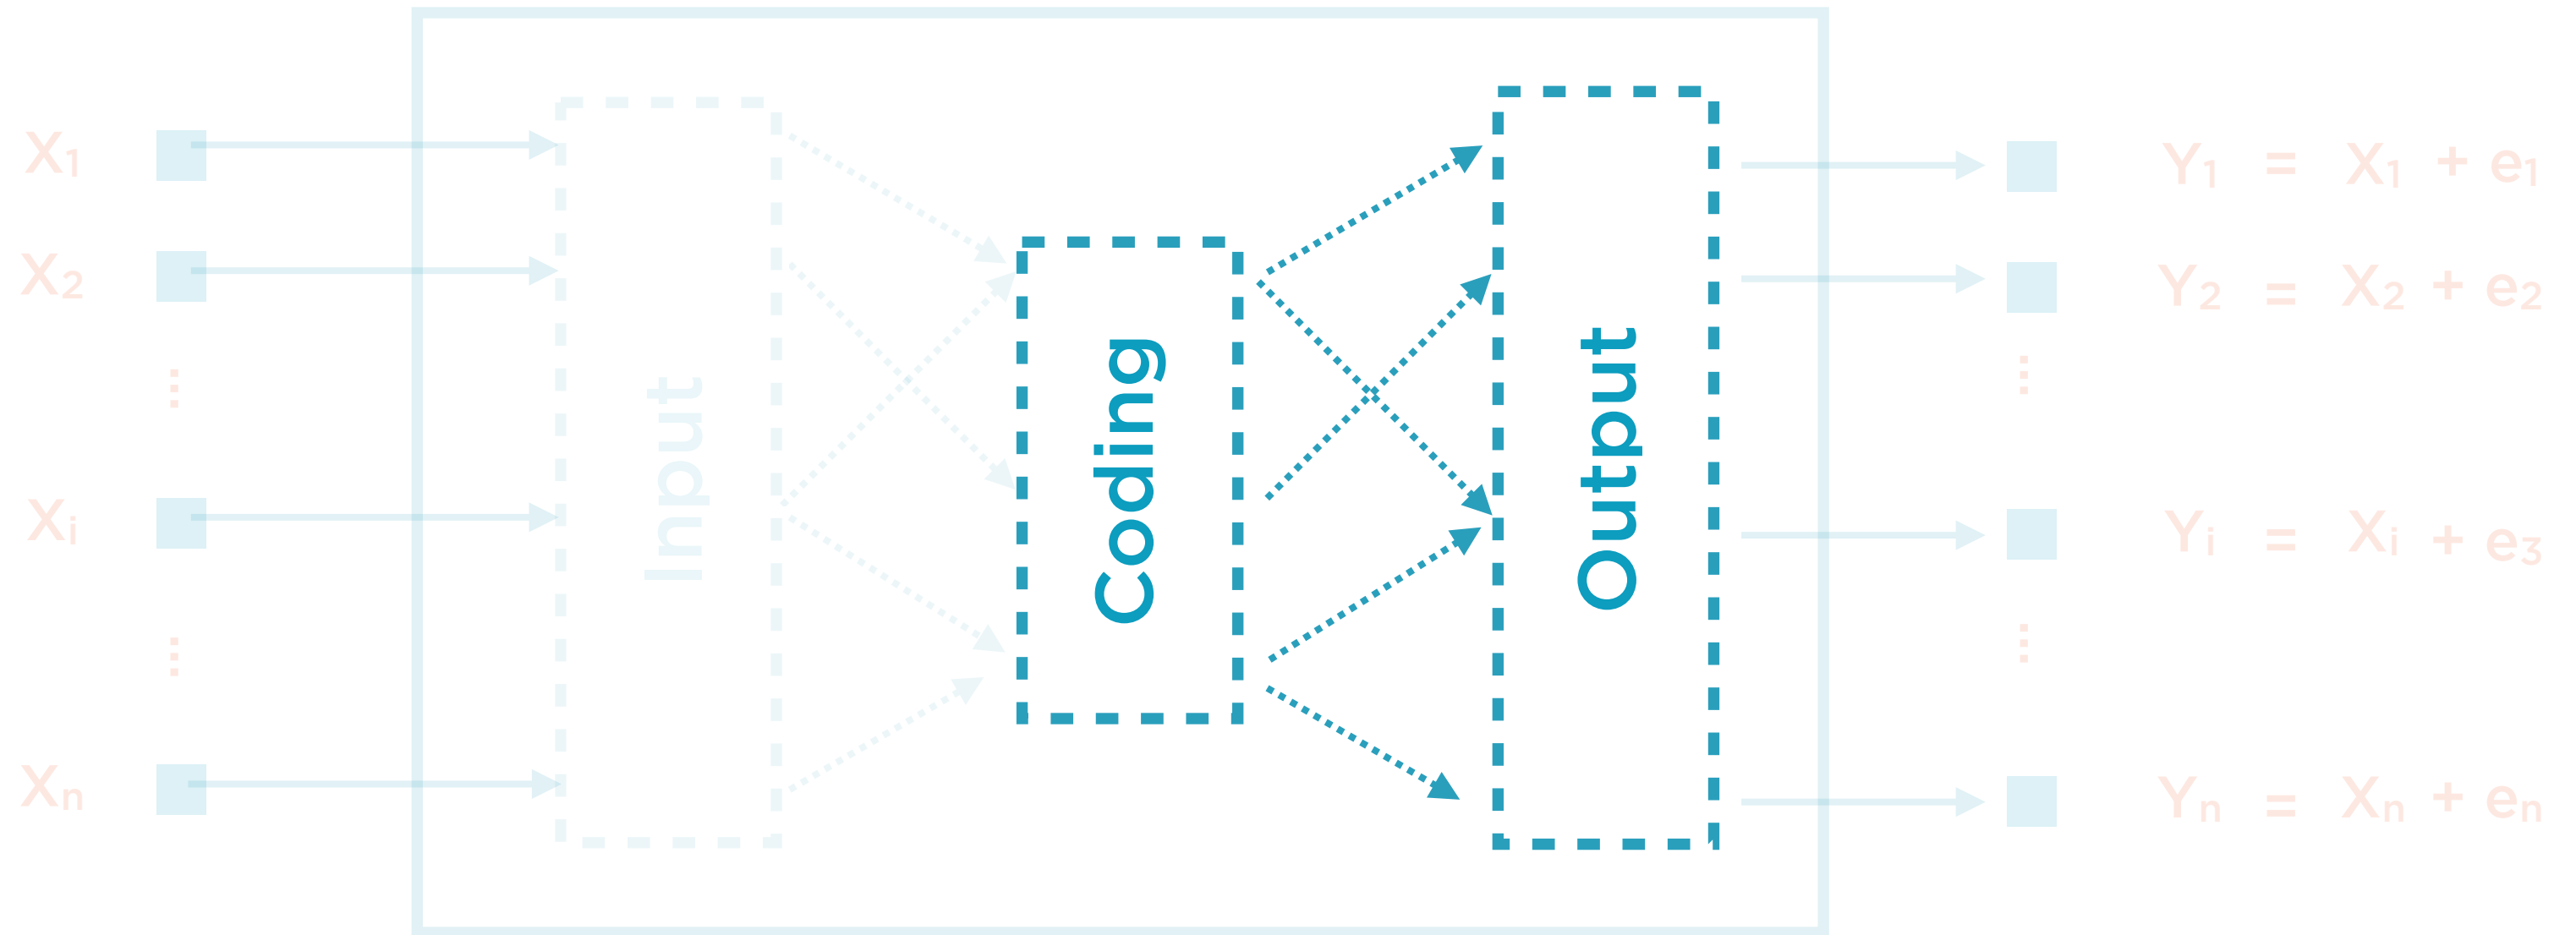
**The central hidden layer is called the coding layer**

# Undercomplete Autoencoder



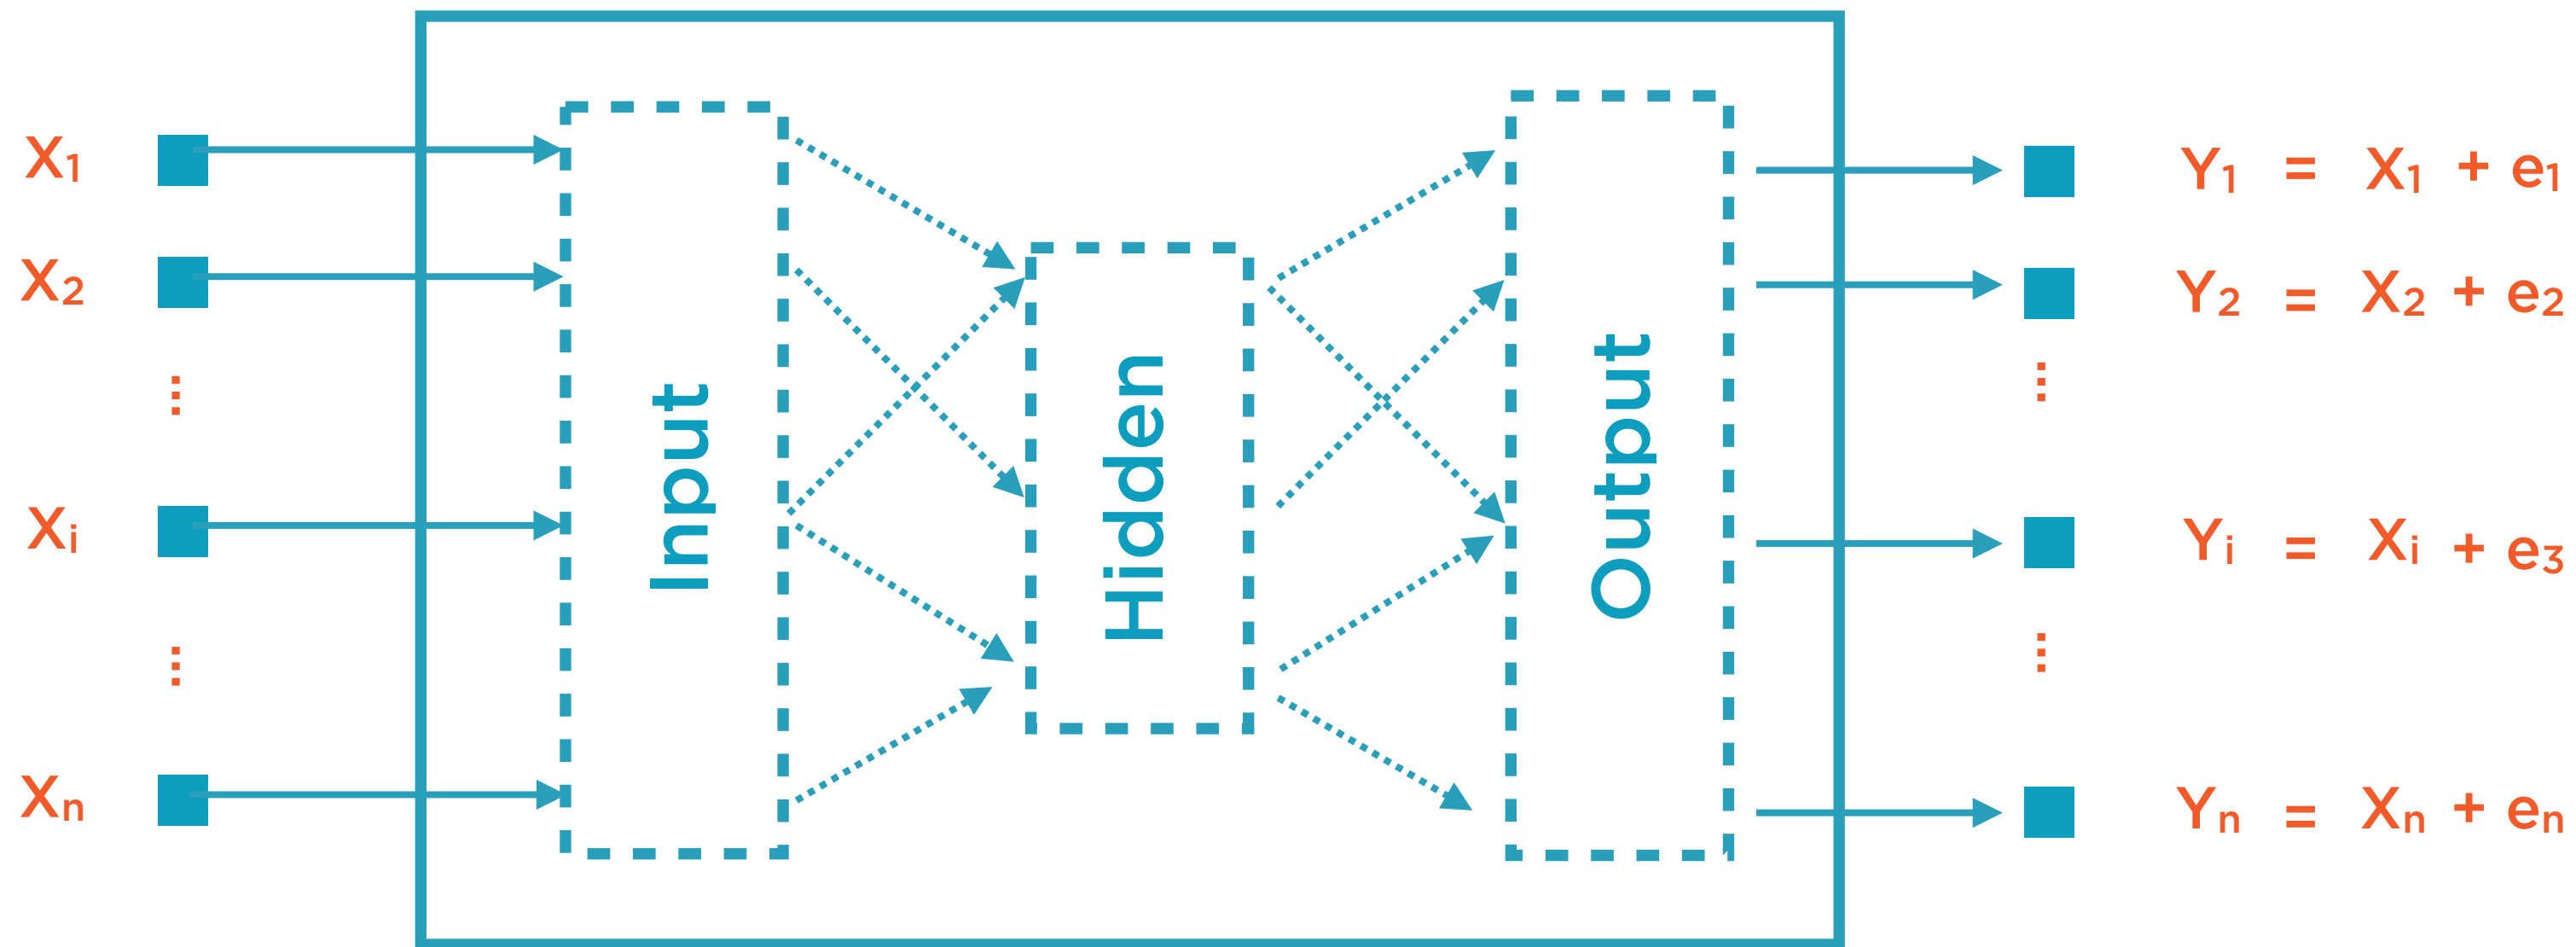
**The first phase - from input to coding - is called encoding**

# Undercomplete Autoencoder



**The second phase - from coding to output - is called decoding**

# Autoencoder





Demo

**Autoencoders to learn latent factors**

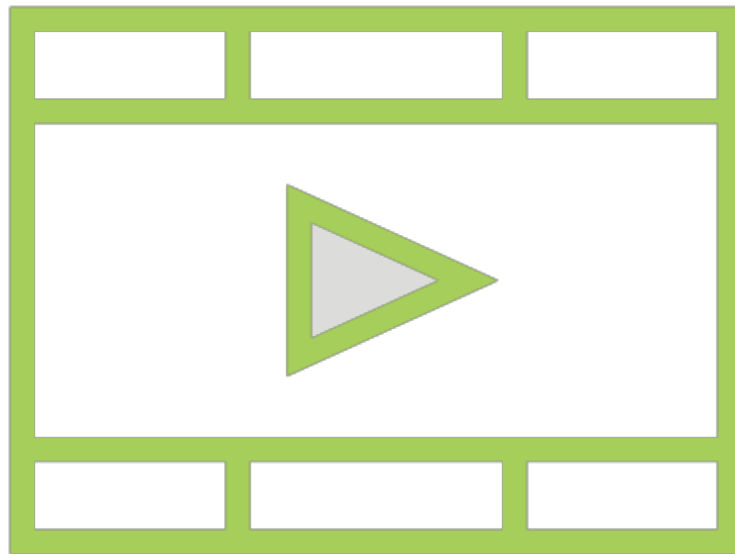
# Summary

**Dictionary learning for dimensionality reduction of image data**

**Feature detection using convolutional layers**

**Autoencoders for dimensionality reduction**

# Related Courses

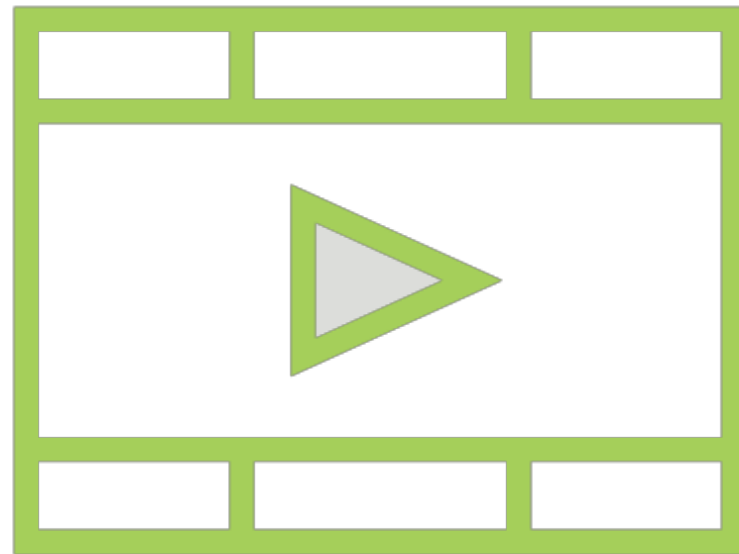


**Building Features from Numeric Data**

**Building Features from Nominal Data**

**Building Features from Text Data**

# Related Courses



**Image Classification with PyTorch**  
**Style Transfer with PyTorch**